

HOW AI WORKS

How AI Memory and Personalization Change Who Gets Recommended

By Scott Tischler · September 25, 2026 · 9 min read

AI recommendations are quietly becoming personal. The same question — "what's the best option for me?" — increasingly returns different answers to different people, because the model now remembers who's asking. ChatGPT and other assistants build a running profile of each user from their history and stated preferences, and they lean on it when they recommend. That means "getting recommended by AI" is no longer one fixed contest with one winner. It's a contest that resolves slightly differently for every person, based on context you don't control. Here's what actually changes, and what you can still influence.

I run Alrecommend.ai, and this is the shift I think most business owners are underestimating. We got used to search returning roughly the same ten links to everyone. Personalized AI recommendations break that assumption. Understanding how they work is the difference between panicking about it and using it.

What personalization actually changes

Let me be precise about the mechanism, because it's easy to overstate. Personalization doesn't rewrite the facts about your business. What the model knows about you as a business — your attributes, your reputation, the consensus of what the web says about you — is still built from public sources and is roughly the same for everyone. What personalization changes is the *filter* applied to that shared knowledge for a specific person.

Think of it in two layers. The bottom layer is the model's general understanding of your category: who the credible options are, what each one is known for, who fits which need. That layer is what classic AEO and entity authority work builds, and it's shared. The top layer is a personal filter: given what this user has told the model and done before, which of those credible options best fits *them*?

Personalization operates almost entirely in that top layer.

The practical consequence is reassuring and uncomfortable at the same time. Reassuring: personalization can't make an unknown business into a recommendation. If you're not in the model's general understanding of your category, no user's profile pulls you in. Uncomfortable: even when you are a credible option, whether you get named for a specific person depends on how well you match their context — and their context is invisible to you.

The three inputs that shape a personalized recommendation

Based on how these systems behave, three things determine which of the credible options a person actually gets recommended.

Stated preferences. When a user has told the assistant something durable — their budget, their location, their priorities, that they prefer premium or prefer thrifty, that they care about a specific feature — the model weights recommendations toward businesses that match. If you are clearly and consistently the "premium, white-glove" option, you win the users whose profiles skew that way and lose the ones optimizing for cheapest. The lesson isn't to be everything. It's to be unmistakably something.

Conversation context. Even without long-term memory, the current conversation personalizes heavily. What the user said three turns ago — the constraints they mentioned, the options they ruled out, the way they framed the problem — shapes the recommendation more than any single ranking factor. This is why being the clean fit for a well-defined need matters more than being a vaguely strong generalist.

Inferred profile. Over time, the assistant infers things the user never explicitly stated — their apparent sophistication, their recurring interests, their patterns. This layer is the least transparent and, honestly, the least worth chasing directly. You can't optimize for a stranger's inferred profile. You can only make sure that whatever segment a user falls into, if you're the right fit for that segment, the model clearly knows it.

Why this is good news for specific, well-defined businesses

Personalization punishes vagueness and rewards clarity, and that's a better world for most honest operators than the one it replaces. In an un-personalized system, the biggest, most generically-known brand tends to be the safe default for everyone. In a personalized system, the model is trying to match a real person to a real fit — which opens the door for the business that is precisely right for a well-defined slice of the market, even if it's not the biggest name overall.

If you serve a specific niche exceptionally well, personalization is your friend. The user whose profile matches your niche is exactly who the model wants to route to you, and your specificity is the signal

that makes the match. The businesses that struggle are the ones trying to be a mild fit for everyone. In a personalized world, a mild fit for everyone is a strong fit for no one, and the model has less reason to surface you for any particular person.

My read: the winning move under personalization is the same move that's always won in positioning — pick a clear "who you're for" and make it legible everywhere. What's new is that the machine now enforces the discipline. Fuzzy positioning used to cost you conversions. Now it costs you the recommendation itself.

What you can and can't influence

You cannot see or shape an individual user's memory, and you shouldn't try. That's their data and their context. What you can do is control the shared bottom layer so thoroughly that whenever a matching user shows up, you're the obvious fit.

Make your "who you're for" unambiguous and consistent across every surface — your site, third-party descriptions, reviews, directories. If the model's general understanding of you is crisp ("the option for X kind of customer who cares about Y"), the personal filter has something clean to match against. If your positioning is mush, personalization has nothing to grab onto and you fall back to being a maybe.

Get corroboration for the specific attributes that define your fit, not just generic praise. "Great service" matches nobody in particular. "Best for first-time buyers who want a lot of hand-holding" matches a specific profile precisely. Reviews and third-party sources that describe your distinctive fit, in the language real customers use, are what let the model connect you to the users you're actually right for.

Answer the qualifying questions on your own pages, too. When a user's profile carries constraints — a budget, a location, a must-have feature — the model is effectively filtering credible options against those constraints. If your site plainly states who you serve, what you cost, where you operate, and what you're genuinely best and worst at, you make yourself easy to match and easy to rule in. Hiding your qualifiers to keep every door open does the opposite: it makes you ambiguous to filter, and an ambiguous option is a skippable one. Honesty about who you're not for is, counterintuitively, one of the strongest ways to win the people you are for.

And resist the urge to broaden. When you feel the pressure to appeal to everyone so you don't miss anyone, remember that personalization rewards the opposite. The narrower and truer your fit, the more decisively you win the people it's meant for.

How to think about measurement when answers vary by person

Personalization breaks the old idea of a single "ranking" you can check once. If the answer differs by user, one test tells you almost nothing. The honest approach is to measure in distributions, not snapshots: test your key prompts across a range of personas and starting contexts — different budgets, different stated priorities, different framings — and look at how often you show up and for which kind of user. That gives you a share-of-answer picture that reflects the personalized reality instead of one lucky or unlucky draw.

Watch especially for the pattern where you win the users who match your niche and lose the ones who don't. That's not a failure — that's personalization working, and it's exactly what you want if your positioning is honest. The failure pattern is showing up weakly for everyone and strongly for no one. If that's what your testing reveals, the fix isn't more visibility work. It's sharper positioning, made consistent everywhere the model can read it.

None of this is a reason to panic. Personalization doesn't hide good, well-positioned businesses from the people they're right for — it does the opposite. It just retires the era of being vaguely known to everyone. The brands that lose sleep over it are usually the ones whose whole strategy was being the biggest generic name in the room; the ones who benefit are the specialists who were always a little too specific for mass search to serve well. Decide who you're for, prove it in the places the model reads, keep it current, and let the machine do the matching.

Key takeaways

- ◆ Personalization changes the filter, not the facts: the model's shared understanding of your business is roughly the same for everyone; the personal profile decides which credible option fits this user.
- ◆ It can't make an unknown business into a recommendation — but among credible options, fit-to-context decides who gets named for a given person.
- ◆ Three inputs shape a personalized recommendation: stated preferences, current conversation context, and the model's inferred profile of the user.
- ◆ Personalization rewards specificity and punishes vagueness — a mild fit for everyone is a strong fit for no one.
- ◆ You can't shape a user's memory; you can make your 'who you're for' crisp and consistent so the personal filter has a clean match to grab.
- ◆ Measure in distributions, not snapshots: test key prompts across personas and contexts to see how often, and for whom, you show up.

Frequently asked questions

Does AI personalization mean my AEO work no longer matters?

The opposite — it makes the foundational work matter more. Personalization operates on top of the model's shared understanding of your category. If you're not in that understanding as a credible, clearly-positioned option, no user's profile pulls you in. AEO and entity-authority work builds that shared layer. Personalization then decides which of the credible options fits a given person. You have to win the base layer before personalization can help you at all.

Can I influence what an individual AI remembers about a user?

No, and you shouldn't try — that's the user's private context. What you can control is the shared, public layer the model builds about your business: your positioning, your attributes, your corroboration across third-party sources. Make it unmistakable who you're the right fit for, in language real customers use, and the personal filter will match the right users to you. Trying to game individual memory is both impossible and the wrong focus.

If answers differ by person, how do I know if I'm winning?

Stop checking a single answer — one test reflects one user's context and tells you little. Test your important prompts across a range of personas: different budgets, priorities, and framings. Then look at the distribution: how often you appear and for which kind of user. Winning the users who match your niche and losing the ones who don't is success, not failure. The warning sign is showing up weakly for everyone and strongly for no one, which points to fuzzy positioning.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.