

ENTITY AUTHORITY

When AI Files You Under the Wrong Category — and Why You Vanish

By Scott Tischler · October 8, 2026 · 9 min read

Before an AI can decide whether to recommend you, it has to decide what you are. It silently files you under a category — "boutique law firm," "SaaS analytics tool," "general contractor" — and that label quietly controls every answer about you from then on. If it files you wrong, you don't show up as a weak answer to your real questions; you don't show up at all, because the model never considered you for them. Fixing a miscategorization is usually higher-leverage than any amount of content, because it's the difference between competing and being invisible.

I spend most of my week at Alrecommend.ai watching how ChatGPT, Gemini, Perplexity, and Google's AI describe the businesses I work with, and category errors are one of the most common and most overlooked problems I find. People assume that if AI "knows" them, it knows what they do. It often doesn't. It knows a fuzzy, inferred version of what you do, and that inference is doing far more work than anyone realizes. This is my read from a lot of observation, not a claim to have seen the inside of any model.

Why the category comes first

Think about how you'd answer "who's a good one of these" for a stranger. You can't even begin until you know what "these" is. AI works the same way. When someone asks for "the best X for Y," the model first forms a picture of the category X, then looks for entities it has already associated with that category, then ranks the ones it can describe. Your name only enters the running if it's attached to the right category in the first place. If the model has you logged as something adjacent but wrong, you're simply not in the pool — and no ranking signal, no review count, no amount of on-page optimization reaches a candidate that was never considered.

This is what makes category errors so nasty. They don't degrade gracefully. A business with a weak reputation still appears occasionally, hedged and low on the list. A miscategorized business appears for the wrong questions and vanishes from the right ones, and the owner has no idea it's happening

because they're testing the queries they *hope* to win, where they were never eligible, instead of noticing they're turning up for queries that bring the wrong customers.

How the model decides what you are

The category isn't something you declare once and have honored. The model infers it, continuously, from the balance of signals it can find. A few things drive that inference more than others.

The words that surround your name most often. If the web repeatedly puts your name near "wedding photography," that becomes your category, even if you've pivoted to commercial work. The model is pattern-matching on co-occurrence, and the loudest, most repeated association usually wins. Your newest and truest positioning can lose to an old one that simply appears more times.

Your own plain-language self-description — if you have one. Models lean on an explicit, clearly stated "we are a [category] that does [specific thing] for [specific customer]" when they can find it. The problem is that most businesses never write that sentence anywhere a machine can lift it. They write aspirational taglines ("reimagining how the world connects") that tell the model nothing about what shelf to put them on, so it falls back on guessing from context.

The company you keep. Directories, "best of" lists, and comparison pages that group you with certain competitors tell the model, in effect, "this one belongs in that set." Being listed alongside the right peers is a strong categorization signal. Being absent from those lists — or worse, listed among the wrong peers — pushes the inference the wrong way.

What you visibly sell and to whom. If half your site talks about one service line and the other half about a newer one you actually care about, the model weights what's most prominent and most corroborated, not what you wish led. I've seen businesses whose homepage buried the thing they're now known for three clicks deep, and the model categorized them by the loud, old thing on the front page.

The ways it goes wrong

Miscategorization isn't one failure; it's several, and they call for different fixes.

The first is the **stale category** — you changed, the web didn't. You used to be a general agency, now you're a specialist, but the body of evidence still describes the old you. The model is accurate about a version of you that no longer exists.

The second is the **too-broad category** — the model has you as "a consultant" or "a marketing company," a bucket so large that you're never the specific answer to anything. You're technically

categorized correctly and still functionally invisible, because real questions carry qualifiers and you're not attached to any of them.

The third is the **adjacent-but-wrong category** — you do fractional CFO work and the model has you as "bookkeeping," or you're a trademark attorney filed under "general legal services." Close enough that it looks fine at a glance, wrong enough that you miss every query that matters and occasionally get surfaced for ones you can't actually serve well.

The fourth, and most damaging for anyone doing something genuinely new, is **no clean category at all**. If you've built something that doesn't fit an existing bucket, the model has nowhere to file you, so it either forces you into the nearest familiar box or treats you as noise. Category creation is powerful for a brand but brutal for machine categorization, because the machine can only recommend you for categories people already ask about.

How to find out where you've been filed

You can't fix this by guessing. You have to go look. Ask each major AI engine open-ended questions about yourself and read the category it reaches for, not just whether it's flattering. Try "what does [your name] do?", "what kind of business is [your name]?", and "what category would you put [your name] in?" Then ask the inverse: "who are the top [the category you believe you're in] for [your customer]?" and see whether you appear. The gap between how you'd describe yourself and the words the model uses back is your miscategorization, laid out in plain sight. Do it across several engines, because they're trained and grounded differently and will often disagree — and that disagreement is itself useful, because it shows which associations are fragile and which are locked in everywhere.

Fixing the label

Repositioning a human audience is a messaging exercise. Repositioning a model is an evidence exercise — you have to change the balance of what the web says, not just what you say. A few moves do most of the work.

- **Write the category sentence, literally, in several places.** "[Name] is a [specific category] serving [specific customer] — we do [specific thing]." Put it on your homepage, your about page, and your profiles, in nearly the same words each time. Give the model a sentence to lift instead of a vibe to decode.
- **Get the right category corroborated off your own site.** One self-description is a claim; the same category echoed in a directory listing, a guest article, a client's write-up, and a review is a fact the model will repeat. Independent agreement is what actually overrides a stale inference.

- **Prune or demote the contradicting signals you control.** If old pages, bios, and listings still describe the thing you no longer do, update or retire them. You're trying to shift the center of gravity, and every stale source is weight on the wrong side.
- **Attach yourself to the right peer set.** Earn your way onto the lists and into the comparisons where your true competitors live. Being named in the right company is one of the fastest ways to move a category inference.
- **If you're category-less, borrow a familiar bucket as the on-ramp.** Anchor yourself to the closest category people actually ask for, win there, and expand the definition over time. The model can't recommend you for a question nobody asks yet.

Why this is worth doing first

I keep coming back to category because it's upstream of everything else people spend their effort on. You can have the best reviews, the cleanest schema, and the most quotable content in your field, and none of it helps if the model has you shelved under the wrong heading — the shoppers for your real category walk right past because, as far as the machine is concerned, you're not in that aisle. Get the label right and all your other work suddenly counts, because it's finally pointed at the questions you were trying to win. It's an unglamorous fix, which is exactly why so few of your competitors have bothered to check. That's your opening.

Key takeaways

- ♦ AI assigns you a category before it decides whether to recommend you — and that inferred label controls which questions you're even eligible for.
- ♦ Miscategorization doesn't degrade gracefully: you don't rank low, you vanish entirely from your real queries while possibly surfacing for wrong ones.
- ♦ The category is inferred from the words that most often surround your name, your own plain-language self-description, the peers you're grouped with, and what you visibly sell.
- ♦ The common failure modes are stale (you changed, the web didn't), too-broad, adjacent-but-wrong, and no clean category at all — each needs a different fix.
- ♦ Find your filing by asking engines open-ended 'what kind of business is this' questions, then the inverse 'who are the top [category]' question to see if you appear.

- ♦ Fixing it is an evidence exercise, not a messaging one: write the category sentence everywhere, get it corroborated off-site, prune contradicting signals, and join the right peer set.

Frequently asked questions

How is being miscategorized different from just having weak AI visibility?

Weak visibility means you're in the running and losing — you appear occasionally, hedged, low on the list. Miscategorization means you were never in the running for your real questions at all, because the model considered a different set of businesses entirely. That's why it's so easy to miss: you keep testing the queries you want to win, where you were never eligible, instead of noticing you're absent from them and sometimes turning up for the wrong ones.

Can I just tell the AI what category I'm in and have it stick?

Not reliably, and not for long. The model infers your category from the balance of evidence across the web, so one self-description competes against everything else that's ever been said about you. If older or louder sources point the other way, your correction gets outvoted. The durable fix is to shift the balance — state the category clearly in several of your own places, then get independent sources to echo the same category so it becomes something the model repeats rather than something you asserted once.

I built something genuinely new that doesn't fit an existing category. What do I do?

This is the hardest case, because a model can only recommend you for categories people already ask about, and a brand-new category has no query volume yet. My honest advice is to anchor to the closest familiar category as an on-ramp — let AI file you somewhere people actually search, win there, and expand the definition of what you do over time. Category creation is great for building a brand with humans, but with machines you usually have to enter through a door that already exists before you can build your own.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.