

HOW AI WORKS

The Company You Keep: How Co-Citation Tells AI You Belong

By Scott Tischler · October 5, 2026 · 9 min read

Here is the direct answer: AI assistants decide who to recommend partly by **the company you keep**. When the web repeatedly mentions you alongside the recognized names in your field — in the same articles, the same roundups, the same expert panels — the model learns, by sheer association, that you belong in that set. This is called co-citation, and it is one of the quietest, most underused levers in getting recommended by ChatGPT, Gemini, and Perplexity. You don't always have to beat the leaders in your category. Sometimes you just have to keep showing up in the same rooms they do.

I think co-citation is badly underrated because it feels indirect. Owners want a single page they can fix to make AI love them. But large language models don't understand your business from one page — they understand it from the web-wide pattern of who gets mentioned with whom. At Alrecommend.ai this is one of the first things I look at when a client is accurate on the web but still never gets named: not what the web says about them, but *who the web says it next to*.

What co-citation actually is

Co-citation is when two names appear together in a source that a third party wrote. Not you linking to yourself, and not you claiming to be a peer of the big players — a neutral article, directory, listicle, podcast show-notes page, or research write-up that names you and names them in the same context. Each time that happens, it is a small, independent vote that you operate in the same space, at a comparable level, for the same kind of buyer.

Models build an internal map of which entities relate to which. If "Scott Tischler" and "Alrecommend.ai" keep appearing beside the established voices on AI search and answer-engine optimization, the model's representation of that topic starts to include me. So when someone asks an unbranded question — "who are the people to follow on getting recommended by AI?" — I am already inside the cluster it reaches for. I didn't earn that by asserting it. I earned it by being genuinely present where that conversation actually happens.

It helps to separate this from a close cousin. **Corroboration** is about the web agreeing on the *facts* about you — same category, same location, same claims. **Co-citation** is about the web placing you in the right *company*. You want both. Corroboration keeps the model from hedging about who you are; co-citation gets you into the shortlist in the first place.

Why AI leans on the company you keep

Put yourself in the model's position. It has to answer "who's good at X?" from an enormous, messy, contradictory web, and it has no way to independently verify quality. So it does what a smart outsider does when they walk into an unfamiliar industry: it looks at who the insiders cluster around. If a dozen independent sources treat a handful of names as the reference points for a topic, that clustering is strong evidence — not proof, but strong evidence — that those names matter. Association is a cheap, scalable proxy for credibility, and models lean on cheap scalable proxies constantly.

This is why a single prestigious mention can do more for your AI visibility than ten pages on your own site. The site pages are you talking about you. The third-party mention beside a recognized name is the market talking about you, and the model weights the two very differently. It is also why co-citation is hard to fake: you can write your own homepage, but you cannot write the sentence in someone else's article that puts your name next to the leader's. That difficulty is exactly what makes it trustworthy to the model.

There's a compounding effect here that's easy to miss. The first time your name appears beside a recognized player, it barely registers. The tenth time, across ten different sources, it becomes part of how the topic is structured in the model's representation. Each mention is weak on its own, but the pattern is strong, and patterns are exactly what these systems are built to detect. This is why co-citation rewards consistency over intensity — a steady presence in the right conversations over a year beats a single splashy mention that never gets repeated.

My honest caveat: co-citation is a signal, not a cheat code. It works when the mentions are real, relevant, and spread across sources the model already respects. It does nothing — or worse than nothing — when it is manufactured in low-quality places, which I'll come back to. And it is a supporting actor, not the lead: if your underlying facts are a mess or your site gives the model nothing clean to quote, strong co-citation can get you named and then let the answer fall apart on the specifics. Build the foundation first, then use co-citation to get you into the conversation.

How to engineer co-citation honestly

The goal is to earn legitimate placements where your name naturally sits beside the names the model already associates with your topic. A few moves that actually work:

Get included in the roundups and "best of" lists that already exist. Someone has probably published "top voices in your space" or "tools to consider." Being added to those lists puts you in the exact co-citation context you want, and many are updated far more often than people assume. A polite, specific outreach — here's who I am, here's the narrow thing I own, here's why I fit — works more often than you'd think.

Show up on the same stages as the recognized names. Guest on the podcasts they guest on. Contribute to the publications they contribute to. Join the expert panels and quote roundups reporters assemble. You are not trying to outshine them in that moment; you are trying to be in the same document, repeatedly, so the association accumulates.

Create the thing worth citing alongside them. Original data, a clear framework, a strong named point of view — these get referenced, and when a writer references your framework next to an established one, you've just manufactured a co-citation the honest way. Owning a specific, nameable idea is the most durable path into the cluster, because it gives other people a reason to mention you when they discuss the category.

Be legible about which cluster you want to be in. If you try to be adjacent to everyone, you end up adjacent to no one. Pick the narrow space you genuinely belong in and concentrate your presence there until the association is unmistakable. Depth beats breadth for co-citation every time.

What backfires

Because co-citation is powerful, people try to shortcut it, and the shortcuts are where it goes wrong. Spraying your name across low-quality directories, paid-link farms, and AI-spun listicles does not build real association — it builds a pattern that the better models increasingly discount or distrust. Worse, it can co-locate you with junk, and guilt by association runs in both directions. If the web keeps naming you beside spam, that is the neighborhood the model files you in.

Buying your way into a fake "as seen in" montage has the same problem: the model is looking at the actual text of real sources, not the logos in your footer. And trying to claim peer status by simply asserting it on your own pages does almost nothing, because self-assertion is the weakest possible signal. The entire value of co-citation comes from it being *other people's* sentences. The moment you manufacture it cheaply, you strip out the thing that made it credible.

How to measure it

You can approximate your co-citation standing without special tools. Ask the assistants directly: "Who are the main people and companies associated with [your topic]?" and see whether you appear and who you appear with. Then run comparison prompts — "How does [you] compare to [recognized

name]?" — and watch whether the model can place you as a peer or stumbles. If it can't locate you relative to the known players at all, you have a co-citation gap, not just a content gap.

Then go look at the sources themselves. Search the key phrases in your space and read who gets named together in the top third-party results. That co-mention map — who appears beside whom — is roughly the map the model has built too. Your job over the next few quarters is to get your name into more of those documents, beside the right company, in places the model trusts. It is slow, it is mostly earned rather than bought, and it is one of the most durable advantages you can build in AI search — because the company you keep is a reputation your competitors can't quietly delete.

Key takeaways

- ◆ Co-citation is when a third-party source names you and a recognized player in the same context — each one is an independent vote that you belong in that set.
- ◆ AI leans on association as a scalable proxy for credibility: if independent sources cluster a handful of names around a topic, the model treats that cluster as the answer to 'who's good at this?'
- ◆ Co-citation is different from corroboration — corroboration gets the web to agree on your facts; co-citation puts you in the right company so you make the shortlist at all.
- ◆ Earn it honestly: get into existing roundups, guest on the same stages as the leaders, publish work worth citing alongside them, and concentrate on one narrow cluster.
- ◆ Faking it backfires — spammy directories and paid 'as seen in' montages either get discounted or file you next to junk, and guilt by association runs both ways.
- ◆ Measure it by asking assistants who's associated with your topic and running peer-comparison prompts, then mapping who gets named together in real third-party sources.

Frequently asked questions

What is co-citation in the context of AI search?

Co-citation is when a neutral, third-party source — an article, roundup, directory, or podcast page — names you and a recognized player in your field in the same context. AI assistants use these repeated co-mentions to learn which businesses belong together in a category, which influences who they name when someone asks for a recommendation.

How is co-citation different from just having good reviews or content?

Reviews and your own content tell AI about you in isolation. Co-citation tells AI where you sit relative to everyone else. You can have flawless reviews and still never be named for an unbranded query because the web has never placed you beside the recognized names. Co-citation is specifically about that placement — the company you keep, in other people's words.

Can I just add myself to directories to build co-citation?

Legitimate, relevant directories and roundups can help, but spraying your name across low-quality or paid-link directories backfires. Better models increasingly discount those sources, and being co-located with spam can file you in the wrong neighborhood. The value of co-citation comes from it being other people's genuine sentences in sources the model trusts — which is exactly what cheap shortcuts strip out.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.