

HOW AI WORKS

How AI Builds Its “Best Of” Lists — and How to Get On Them

By Scott Tischler · September 18, 2026 · 8 min read

When a customer asks ChatGPT, Perplexity, or Google's AI for “the best project management tool for a small agency” or “the best physical therapist in Austin,” the model is not reading a leaderboard off a shelf. There is no master ranking sitting in a database that it looks up and reads back. It builds the list in the moment, out of whatever names it can pull together and defend, and then it orders them by how confidently it can say something specific about each one. If you understand that the list is *assembled* rather than *retrieved*, the whole game of getting onto it changes.

I spend most of my week watching these answers form and re-form for the businesses I work with at Alrecommend.ai, and the pattern is consistent enough that I'll lay it out plainly. This is my read from a lot of observation, not a claim to have seen inside anyone's model.

A “best of” answer is manufactured on the spot

Here is the sequence, as best I can reconstruct it. The model takes the query and forms an internal picture of what “best” means for that specific ask — the category, the constraints (small agency, Austin, budget, beginner), and the shape of a good answer. Then it gathers candidate names: some from what it absorbed during training, some from live retrieval if the tool is browsing. Then it does the part everyone underestimates. For each candidate, it asks itself a quiet question: **can I say something specific and defensible about why this one belongs here?** The candidates it can back up survive. The ones it can only name vaguely get dropped, because a model that lists a name it can't justify is a model that's about to hallucinate, and the better systems are tuned to avoid exactly that.

That last step is the whole ballgame. Being *known* to the model gets you considered. Being *describable* — in concrete, corroborated terms — is what gets you kept. Most businesses that feel invisible in AI answers are actually in the candidate pool. They just get cut at the descriptability step because nothing the model can find lets it finish the sentence “they're a good choice because...”

Three things decide whether you make the cut

When I watch which names survive and which get dropped, it comes down to three signals working together.

Corroboration. The model wants to see the same claim about you in more than one independent place. If your site says you specialize in Shopify migrations and nothing else on the web echoes that, the claim is fragile. If three review platforms, a directory, a podcast transcript, and two client case studies all reinforce it, the model can lean on it hard. One source is a rumor. Several agreeing sources become a fact the model will repeat. This is why earned mentions on other people's properties often move the needle more than another page on your own site.

Structure. “Best of” queries reward businesses whose information comes in a shape the model can lift cleanly — a clear category label, an explicit specialty, plainly stated qualifiers (who it's for, who it's not for, price band, location). A model assembling a shortlist is looking for a tidy phrase it can attach to your name. If you've written that phrase somewhere in plain language, you've done its work for it. If it has to infer your specialty from vibes, it often won't bother.

Recency and aliveness. For any list where things change — tools, providers, local services — the model leans toward names it has fresh evidence still exist and still do the thing. A business that stopped publishing three years ago reads as a risk. Signs of current activity (recent content, recent reviews, updated listings) tell the model you're safe to recommend today, not just that you existed once.

Where your candidates actually come from

It helps to know where the names in the pool originate, because it tells you where to invest. Some candidates come from the model's training — things it absorbed about your category before its knowledge cutoff. You can influence that pool over time by building a durable, widely-referenced presence, but you can't change it this quarter and you can't change what a given model already learned. The other source is live retrieval: when the tool browses or searches at answer time, it pulls in current candidates and evidence. That pool you *can* influence quickly, because it reflects what's findable and corroborated on the web right now.

This split matters for how you set expectations. If you're a newer business, you may barely exist in training memory yet, and pushing on that is slow, frustrating work. But you can be very present in the retrieval pool almost immediately by making sure the live web — your listings, reviews, mentions, and current content — says clearly and consistently what you do. In practice the retrieval path is where most of the winnable ground is for the businesses I work with, and it's the one people underinvest in because it feels less permanent. It's less permanent, but it's also the one that pays this month.

You have to match the qualifier, not just the category

Here's the mistake I see most often. A business optimizes to be “the best marketing agency” and wonders why it never appears. But almost nobody asks AI for “the best marketing agency” in the abstract — the answers to that are dominated by the giants and it's a near-useless query anyway. Real people ask for “the best marketing agency *for dental practices*” or “*for early-stage SaaS*” or “*that does HIPAA-compliant email*.” The qualifier is where the shortlist is actually decided, because the qualifier is where the field narrows to a size where a specific business can win.

So the winnable question isn't “how do I rank for my category.” It's “which qualified versions of my category am I the obvious, corroborated answer to?” Pick the qualifiers you can genuinely own, then make sure the web says, in several places, that you own them. Owning “best bookkeeper for restaurants in Ohio” is achievable in a way that “best bookkeeper” never will be, and it's also the query that brings you a customer who actually fits.

The playbook for earning a slot

None of this requires gaming anything. It requires making yourself easy to include. Concretely, this is what I'd do.

- **Write your one-line specialty in plain words, everywhere.** The exact phrase you want the model to attach to your name should appear on your homepage, your profiles, and ideally in how others describe you. Don't make it guess.
- **Get the same specialty claim corroborated off-site.** Reviews that mention what you're good at, a directory listing with the right category, a guest article, a client quote on someone else's site. Independent agreement is what converts a claim into a fact the model repeats.
- **Target qualified queries, not the naked category.** Build a page and a body of proof for each qualifier-plus-category you can honestly win. That's where the shortlist has room.
- **Show you're alive.** Keep something current — recent reviews, a refreshed listing, new content. Freshness is a tiebreaker and a safety signal at once.
- **Remove contradictions.** If different sources describe you differently, the model hedges. Make the web tell one consistent story about who you are and who you serve.

What I'd actually measure

Don't measure whether you “rank.” Measure inclusion and describability. Take the ten or fifteen qualified “best of” questions a real customer would ask, run them across ChatGPT, Perplexity, and Google's AI, and log two things: are you named at all, and when you're named, is the reason the model

gives *accurate and specific*? A wrong or generic reason is almost worse than being skipped, because it means the model has picked up a fuzzy or mistaken version of you and is now repeating it to buyers. Fixing the descriptability — giving the web a clean, corroborated sentence about what you're best at — is usually the highest-leverage work available, and it's the same work whether the engine is ChatGPT today or something we haven't named yet.

The businesses that win the “best of” era won't be the loudest. They'll be the most *quotable* — the ones that made it effortless for a machine assembling a shortlist to include them and say exactly why. That's a reputation problem dressed up as a technology problem, which is good news, because reputation is something you can actually build on purpose.

Key takeaways

- ♦ AI “best of” lists are assembled live from candidates the model can name AND justify — they are not read off a fixed ranking.
- ♦ Being known to the model gets you considered; being describable in specific, corroborated terms is what keeps you on the list.
- ♦ Three signals decide inclusion: independent corroboration of your specialty, clean structure the model can lift, and evidence you're currently active.
- ♦ The shortlist is won at the qualifier (“for dental practices,” “in Austin”), not at the bare category — own the qualified versions you can defend.
- ♦ Contradictory descriptions across the web make the model hedge; one consistent story about who you serve is worth more than another page on your own site.
- ♦ Measure inclusion and whether the reason the model gives for you is accurate, not whether you “rank.”

Frequently asked questions

Why does AI recommend competitors instead of me when I'm objectively better?

Usually because the competitor is easier to describe, not because a model judged quality. If several independent sources say what a competitor specializes in and only your own site says what you specialize in, the model can back up their inclusion and can't confidently back up yours — so it keeps them and drops you. The fix is corroboration: get the web to echo, in multiple independent places, the specific thing you're best at.

Do I need to appear on listicle-style “best of” web pages to show up in AI's best-of answers?

It helps but it isn't required. Those pages are strong corroboration because they're independent third parties naming you in a category, so they feed the model exactly the kind of evidence it wants. But the underlying need is agreement across sources — reviews, directories, mentions, case studies — and a human-written listicle is just one convenient form of that.

Should I try to be the “best” at my whole category or a narrow niche?

A narrow, honest niche, almost always. Real queries carry qualifiers, and the qualifier is where the field narrows enough for a specific business to win. Owning “best X for [specific customer] in [place]” is both more achievable and more valuable, because it also brings you the customer who actually fits what you do.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.