

HOW AI WORKS

How AI Actually Finds Its Answers: RAG, Grounding, and Who Gets Cited

By Scott Tischler · September 18, 2026 · 9 min read

The direct answer: When you ask ChatGPT, Gemini or Perplexity to recommend a business, most of the time it isn't answering from memory. It runs a search, pulls a handful of documents, reads them, and writes an answer grounded in what those documents say. That process is called retrieval-augmented generation — RAG — and understanding it is the closest thing there is to seeing the machine that decides who gets cited. If you know how retrieval works, you know where the leverage is.

I'm going to explain RAG the way I'd explain it to a business owner, not an engineer, because the mechanics have direct consequences for whether AI names you. Skip the jargon and the shape is simple: the model is only as good as the documents it retrieves, and you have more influence over those documents than you think.

What actually happens between your question and the answer

Picture the assistant as a very fast research assistant rather than an oracle. When your question arrives, it doesn't reach into a fixed memory of the world. It does something more like what a person would do who wants to sound informed and current.

First it **interprets your question** and turns it into one or more searches. "Who's a good commercial real estate broker in Dallas" might become several queries the system runs behind the scenes. Then it **retrieves** — it fetches a set of candidate documents from a search index or its own store: web pages, business listings, articles, reviews, structured data. Then it **reads and ranks** those documents, deciding which passages are relevant and trustworthy enough to use. Finally it **generates** the answer, composing a fluent response that leans on the retrieved passages and, increasingly, cites them.

The critical thing to internalize: the model's built-in knowledge is the fallback, and retrieval is the foreground. If good, fresh documents about you exist and are retrievable, they shape the answer. If they don't, the model falls back on whatever vague impression it absorbed during training — which for

most local and mid-market businesses is thin, outdated, or nonexistent. That gap between "retrievable now" and "vaguely remembered" is where visibility is won and lost.

Why retrieval changes the whole optimization game

For twenty years, SEO was about ranking a page so a human would click it. RAG keeps some of that — being findable still matters — but it changes what "winning" means. The model isn't trying to get you a click. It's trying to *extract facts it can safely repeat*. Those are different jobs, and they reward different things.

Under RAG, three properties suddenly matter a lot more than they used to.

Retrievability. If your information isn't in the indexes these systems search, or your site blocks their crawlers, you're invisible to retrieval no matter how good your business is. You can't be cited from a document that was never fetched. This is the unglamorous foundation — accessible pages, clean structure, presence in the listings and sources the engines actually read.

Extractability. Retrieval pulls passages, not whole pages. A page that buries its key facts in clever marketing prose gives the model nothing clean to lift. A page that states plainly what you do, where, for whom, and what makes you different hands the model quotable material. I've watched a great business lose to a mediocre one purely because the mediocre one wrote in extractable sentences and the great one wrote in a fog of adjectives.

Corroboration. RAG rarely rests on a single source. When several independent documents agree about you — your own site, a directory, an article, a review corpus all telling the same story — the model treats the fact as safe and states it confidently. When sources contradict each other, the model hedges or drops the claim rather than risk being wrong. Consistency across the web is, functionally, a retrieval optimization.

The grounding decision: memory versus search

One subtlety worth understanding is that the assistant is constantly making a judgment call about whether to trust its own memory or go retrieve. This is sometimes called grounding. For a stable, well-known fact — the capital of a country — it may answer straight from training. For anything current, specific, or local — which business to recommend right now — it leans toward retrieval, because its training data is frozen at a cutoff date and it "knows" that recommendations go stale.

This matters for you in two directions. On one hand, it's good news: even if the model learned nothing about you during training, live retrieval gives you a path to be surfaced today, without waiting to be absorbed into some future model. On the other hand, it means retrievable, current, consistent

information is doing the heavy lifting on your behalf every single time. If that information is stale or contradictory, retrieval will faithfully surface the stale, contradictory version of you.

My read, not a guarantee: as these systems mature, they're leaning *more* on retrieval and grounding, not less, precisely because it's how they stay current and cite sources. That trend favors businesses that keep their web presence clean, fresh and consistent — and it quietly punishes the ones who set up a website once in 2019 and never touched it again.

The mistakes that quietly make you unretrievable

Because retrievability is the foundation, it's worth naming the self-inflicted wounds I see most often — the ones that take a good business out of the running before the model ever gets to judge it.

The first is **blocking the crawlers you want to be cited by**. A lot of sites tightened their robots rules to keep AI bots out during the scraping panic, then forgot they did it. If you've told the engines not to read you, they will honor that and quote a competitor instead. Retrieval can't fetch a door you locked.

The second is **hiding your facts inside images, PDFs, or JavaScript that renders late**. If the plain text of who-what-where isn't easily readable in the page itself, retrieval may never extract it. Beautiful design that a machine can't parse is, for these purposes, an empty page.

The third is **letting your identity contradict itself across the web**. A different business name here, an old address there, a service you dropped two years ago still listed somewhere else. Every contradiction gives the model a reason to hedge. Retrieval surfaces the mess, and a hedging model is a model that recommends someone clearer.

None of these are exotic. They're the boring hygiene that decides whether the more interesting authority work even has a chance to matter.

The engines retrieve differently — but the fundamentals hold

People want to know whether they should optimize differently for ChatGPT versus Perplexity versus Google's AI answers. The honest answer is that the engines do differ in the details, but the fundamentals are shared, and chasing the details before you've nailed the fundamentals is a waste.

Perplexity is the most openly retrieval-first — it's built around searching, citing, and showing its sources, so extractable, citable content is rewarded almost immediately. ChatGPT blends its own reasoning with live browsing and increasingly names sources, so it rewards both a clean footprint and genuine authority the model has absorbed. Google's AI answers sit on top of the search index Google has always run, so classic findability and structured data carry extra weight there. Gemini leans on Google's grounding stack for the same reason.

But notice what every one of those has in common: they all fetch documents and prefer sources that are accessible, clearly written, corroborated, and current. Optimize for those four properties and you're feeding all of them at once. The engine-specific nuances are real, but they're refinements on a shared foundation — not four separate games. My read: build the foundation first, and worry about the per-engine tuning only once you're consistently getting retrieved at all.

What this means you should actually do

You don't need to understand embeddings or vector databases to act on RAG. You need to make yourself the easiest, most trustworthy source for a machine to retrieve and quote. In practice that comes down to a few things.

Be retrievable. Make sure the pages that describe you are accessible to AI crawlers, well-structured, and present in the directories and platforms the engines pull from. Don't accidentally block the very systems you want to be cited by.

Be extractable. State your core facts in plain, self-contained sentences — who you are, what you do, where, for whom, what's distinctive. Answer the real questions customers ask, in language a model can lift verbatim without having to interpret your cleverness.

Be corroborated. Get the same accurate story told about you in more than one place. Third-party mentions, consistent listings, and a coherent identity across the web are what turn a retrievable fact into a confident recommendation. Contradiction is the enemy; when the web disagrees about you, the model hedges.

Stay fresh. Because retrieval favors current sources, information that's obviously maintained gets pulled over information that looks abandoned. Keeping your facts up to date isn't vanity — it's feeding the retrieval layer the version of you that you actually want repeated.

Here's the reframe I'd leave you with. Stop thinking of AI as a judge that has already formed an opinion of your business. Most of the time it's a researcher forming an opinion in real time, from whatever it can retrieve in the moment you're asked about. Your job isn't to change its mind. Your job is to make sure that when it goes looking, the clearest, most consistent, most current answer it finds is the truth about you.

Key takeaways

- ♦ Most AI recommendations aren't answered from memory — the assistant runs a live search, retrieves documents, reads them, and grounds its answer in what it pulls. That process is RAG (retrieval-augmented generation).

- ♦ The model is only as good as the documents it retrieves. Its built-in training knowledge is the fallback; retrievable, current sources are the foreground — and that's where you have leverage.
- ♦ RAG rewards three properties: retrievability (can it fetch you), extractability (can it lift clean facts), and corroboration (do multiple sources agree). All three beat old-style click-chasing SEO.
- ♦ Assistants constantly decide between trusting memory and searching. For anything current or local, they lean toward retrieval — which means you can be surfaced today without waiting to be baked into a future model.
- ♦ Contradiction is the enemy: when web sources disagree about you, the model hedges or drops the claim. Consistency across the web functions as a retrieval optimization.
- ♦ Reframe AI as a researcher forming an opinion in real time, not a judge who already decided. Your job is to make the clearest, most consistent, most current source it finds be the truth about you.

Frequently asked questions

What is retrieval-augmented generation (RAG) in plain terms?

It's the way modern AI assistants answer questions using live information instead of only their training. When you ask something, the system turns your question into searches, fetches a set of candidate documents, reads and ranks them, then writes an answer grounded in those passages — often citing them. In short: it researches before it answers, and the documents it retrieves shape what it says about you.

If AI searches the web to answer, do I still need to worry about what the model 'knows'?

Less than you'd think, and that's good news. Because assistants lean on retrieval for anything current or local, you don't have to wait to be absorbed into a future version of the model to get recommended. If clear, consistent, up-to-date information about you is retrievable right now, it can shape today's answer. The flip side: if that information is stale or contradictory, retrieval will faithfully surface the stale, contradictory version of you.

How do I make my business easy for RAG to cite?

Make yourself the easiest, most trustworthy source to retrieve and quote. Be retrievable — accessible pages, clean structure, presence in the listings the engines actually read, and don't block AI crawlers. Be extractable — state your core facts in plain, self-contained sentences a model can lift verbatim. Be corroborated — get the same accurate story told in more than one place so the model treats it as safe. And keep it fresh, because retrieval favors sources that look maintained over ones that look abandoned.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express,

MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.