

HOW AI WORKS

Mentioned vs. Recommended: Why AI Names You but Doesn't Pick You

By Scott Tischler · October 3, 2026 · 8 min read

Short answer: being *mentioned* by an AI assistant and being *recommended* by it are two different events, and most businesses are winning the first while losing the second. A mention means the model knows you exist and will name you when asked for a list. A recommendation means the model puts you forward as the answer — the one it would tell a friend to call. The gap between those two is where almost all of the real commercial value sits, and closing it is a different job than getting mentioned in the first place.

I spend my days looking at how ChatGPT, Gemini, Perplexity and Google's AI answers actually name businesses, and the single most common reaction I get from owners is some version of "wait, I *am* showing up." Yes — in a list of six. Buried in the middle. With a neutral one-line description that could belong to any of your competitors. That is a mention. It feels like a win and it rarely converts like one.

Why a mention and a recommendation are not the same thing

When you ask an assistant an open question — "who are the AEO consultants worth knowing" — it tends to produce a roster. Rosters are cheap for a model to generate, because it just has to recall entities that are associated with the topic. Association is a low bar. If your name co-occurs with the subject often enough across the data the model learned from, you'll surface. That's a mention.

A recommendation is a higher-confidence act. Now the model isn't just recalling that you're associated with the category — it's making a judgment that you are a **safe, specific, defensible pick** for this particular person's situation. Models are, for good reason, cautious about staking a recommendation on a business they only vaguely understand. The downside of a bad recommendation is larger than the downside of a bland list, so the bar rises. The businesses that clear it are the ones the model can describe with specifics, place in a context, and back up with corroboration it has seen in more than one place.

Here's the way I think about it: a mention proves you're *known*. A recommendation proves you're *understood*. Most AEO advice gets you known. Very little of it gets you understood.

There's a second thing going on, too. When the question is vague, the model has room to be generous and hand out a long list. When the question gets specific — a real person with a real constraint — the list collapses to one or two names, because now the model is effectively being asked to take a position. Vague questions reward mentions. Specific questions reward recommendations. And buyers who are close to spending money almost always ask the specific version.

The quiet cost of living in the list

It's worth being blunt about why this matters, because "I show up" feels like success and the data underneath it usually isn't. A mention in a six-item list is seen, skimmed, and mostly forgotten. The person asking didn't want a directory; they wanted a decision made easier. If you're one neutral line among six, you've given them nothing to decide with, and they do what people do with lists — they pick the name they already recognize, or the one with a reason attached, and that name isn't necessarily you.

The deeper cost is that mentions are unstable. Because a list is low-stakes for the model, the contents of that list shift from run to run and model to model. You're in it today, out of it tomorrow, and you have no idea why, because you never earned a durable reason to be there. Recommendations are stickier, precisely because they're built on specifics and corroboration that don't evaporate between sessions. Chasing mentions is chasing a moving target; earning recommendations is building something that holds.

The signals that turn a mention into a pick

When I look at businesses that graduate from "named in the list" to "put forward as the answer," a few things are almost always true. None of them are tricks. They're the things that give a cautious model enough to commit.

You're described with specifics, not adjectives

If the only thing the web says about you is that you're "a trusted provider of quality service," the model has nothing to differentiate you with, so it treats you as interchangeable — list material. The businesses that get recommended are the ones with a crisp, repeated, specific description: who they're for, what they actually do that's different, the niche they own. When the model can say *why* you and not just *that* you, it will.

The same story shows up in more than one place

A claim that only appears on your own website is a claim the model has to take on faith, and it discounts faith. The same claim appearing on your site, in a directory, in a podcast transcript, in a third-party article — that's corroboration, and corroboration is what lets a model move from "they say this about themselves" to "this is true." Recommendations ride on corroborated specifics. Mentions survive on self-description.

You map cleanly to a real intent

Recommendations happen inside a context: a city, a budget, a use case, a type of customer. If the model can only associate you with a broad category, you'll get listed under the broad query and dropped the moment the question gets specific. If it can associate you with the *specific* situation a buyer is actually in — "best X for a small B2B team," "Y studio for beginners over 40" — you become the natural pick when that situation comes up. The narrower and more real the intent you own, the more often you're the recommendation rather than one line in the list.

Nothing in the picture is contradicting you

A model that finds conflicting signals — a stale address, a description that fights with your reviews, a claim it can't reconcile — doesn't usually call you out. It just quietly declines to recommend you and reaches for the business it's more sure about. Cleanliness and consistency don't win on their own, but their absence is enough to lose. Being the safe pick means there's nothing the model has to hedge around.

How to tell which one you're getting

You cannot manage this by vibes, and you cannot manage it from your own logged-in account, where personalization muddies the result. The practical move is to test the questions a buyer would actually ask, in a clean session, more than once, across more than one assistant — and then read the *shape* of the answer, not just whether your name appears.

Ask yourself: Did the model volunteer me without being prompted with my name? Did it put me near the top, or bury me? Did it describe me with something specific, or with a generic line? Did it hedge ("you might also consider...") or commit ("a strong fit would be...")? When it got more specific about the buyer's situation, did I survive the narrowing or fall out of it? Those distinctions are the difference between a mention and a recommendation, and they're visible if you look for them instead of just scanning for your name.

One honest caveat: answers vary between runs, and a single test is noise. I never draw conclusions from one prompt on one day. Run the set a handful of times, look at how often you're the pick versus the list filler, and track that over weeks. The pattern is the signal; the single answer isn't.

The work that closes the gap

If you're getting mentioned but not recommended, the fix is almost never "more content." It's sharper, more specific, more corroborated content about a narrower thing. Pick the intents you actually deserve to win — the specific situations where you genuinely are the right call — and make the case for them clearly, in your own words, on your own site, and then get that same case echoed somewhere a model trusts more than you: earned coverage, a credible directory, a guest appearance, an interview, original data others cite.

The goal is to hand the model everything a cautious recommender needs: a specific description it keeps seeing, a context you clearly fit, corroboration from more than one source, and no contradictions to trip over. Do that for a handful of real buyer intents and you stop being the name in the middle of the list. You become the one the assistant puts forward when someone asks what they should actually do.

My read, not a guarantee: the businesses that win the next few years of AI search won't be the ones mentioned most. They'll be the ones the models are most confident recommending. Those are not the same race, and most of your competitors are still running the wrong one.

Key takeaways

- ♦ A mention means an AI knows you exist and will list you; a recommendation means it puts you forward as the answer. Most businesses win the first and lose the second.
- ♦ Recommendations clear a higher confidence bar because models are cautious about staking a pick on a business they only vaguely understand.
- ♦ Specific, repeated descriptions beat generic adjectives — the model needs to know why you, not just that you.
- ♦ Corroboration across multiple sources is what lets a model move from 'they claim this' to 'this is true,' which recommendations depend on.
- ♦ Owning a narrow, real buyer intent makes you the natural pick when that exact situation comes up; broad association only earns list filler.
- ♦ Test in clean sessions across multiple assistants, read the shape of the answer, and track the pattern over time — a single run is noise.

Frequently asked questions

Is being mentioned by AI still valuable even if I'm not recommended?

It's a necessary first step — you can't be recommended if the model doesn't know you exist. But a mention in a long list rarely converts, because the buyer is looking for a specific pick, not a roster. Treat a mention as proof you're known and then do the work to become understood: sharper specifics and outside corroboration for the narrow intents you actually deserve to win.

Why does AI list my competitor as the recommendation instead of me?

Usually because the model can describe them more specifically and back that description up in more than one place. If your competitor owns a clear, repeated, corroborated story about a particular kind of customer or use case, and you're described in generic terms, the model treats them as the safer, more defensible pick. It's less about who's better and more about who the model understands with confidence.

How do I know if I'm getting mentioned or recommended?

Test the real questions a buyer would ask, in a clean (logged-out) session, several times, across ChatGPT, Gemini and Perplexity. Then look at the shape of the answer: did the model volunteer you near the top with a specific description and a committed tone, or bury you mid-list with a generic line? Did you survive when the question got more specific? The pattern across runs tells you which one you're getting.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.