

PLAYBOOK

The Modifier Problem: How AI Routes “Best,” “Cheapest,” and “Luxury” Queries

By Scott Tischler · September 23, 2026 · 8 min read

When someone asks an AI assistant for a recommendation, they almost never ask a neutral question. They ask for the *best* dermatologist, the *cheapest* CRM, a *luxury* wedding venue, a good option *near me*, the *fastest* way to get a passport photo. That single adjective — the modifier — does more to decide who gets named than the noun ever does. So if you want to be the business an AI recommends, stop optimizing for "CRM" or "wedding venue." Optimize for a specific modifier. The model resolves each one differently, using different evidence, and no honest business can win all of them at once.

The modifier is the real query

"Best plumber," "affordable plumber," "emergency plumber," and "family-owned plumber" are four different questions that route to four different answers — even though the noun never changes. The noun sets the category. The modifier sets the ranking function. It tells the model what to weigh and, just as importantly, what to ignore.

Here's the mistake I see constantly: businesses write a generic page that answers the noun and quietly assumes the modifier will take care of itself. It won't. A page that says "we're a full-service plumbing company serving the metro area" answers no modifier at all. It's legible as a plumber and invisible as the *right* plumber for any qualified question. So the business shows up as a name in the pile and never as the answer to the query the buyer actually typed.

And buyers are typing modifiers more than ever, because conversational AI invites them to. When search was a keyword box, people typed "plumber near me" and stopped. When the interface is a chat, they say "I need an affordable plumber who can come today and won't upsell me." Three modifiers in one breath. The natural-language interface pulls the hidden qualifier out into the open, which means the businesses that win are the ones whose evidence matches those qualifiers.

How AI actually resolves a modifier

I want to be careful here: this is my read from watching how these systems behave across thousands of prompts, not a claim about their internal machinery. But the pattern is consistent enough that I'd bet on it.

The model maps each modifier to a set of proxy signals it can actually find in text, then favors sources that make those proxies explicit and — this part matters — corroborated by someone other than you. "Best" leans on consensus: awards, ranked lists, and praise repeated across independent sources. "Cheapest" and "affordable" lean on explicit price data and comparative language. "Luxury" and "premium" lean on positioning language, the *height* of your prices, materials and craftsmanship detail, and the company you're mentioned alongside. "Fastest" and "same-day" lean on stated turnaround times and reviews that specifically mention speed. "Near me" leans on place associations and local corroboration.

The unifying idea is this: the model isn't measuring reality. It's measuring how legible your reality is in text. It cannot verify that you're the best; it can only see whether independent sources describe you that way. It cannot confirm you're fast; it can only find where someone said so. Whichever modifier your evidence makes legible is the modifier you'll win — and if your evidence is mush, you win nothing.

The five modifier families and what each one rewards

Almost every buyer modifier falls into one of five families, and each rewards a different kind of proof:

- **Superlative** (best, top, leading, #1): rewards independent consensus. You cannot self-declare your way into this one — the model discounts your own claims and looks for third parties who agree.
- **Price** (cheapest, affordable, budget, value, cost-effective): rewards explicit, current numbers and honest comparison. Vague "competitive pricing" language earns you nothing here; a real price does.
- **Prestige** (luxury, premium, high-end, bespoke, boutique): rewards positioning, price height, craftsmanship detail, and adjacency — being named in the same breath as other premium options.
- **Speed and convenience** (fastest, same-day, 24/7, instant, easy): rewards specific, verifiable operational claims and reviews that echo them.
- **Fit and segment** (best for small teams, for beginners, for enterprises, for families, for runners): rewards narrow, explicit "who this is for" language. For most businesses, this is the winnable one.

Pick the modifier you can actually win

Most businesses pour their energy into "best," which is the single hardest modifier to earn because it depends on consensus you don't control. You can spend a year producing content and still lose "best

CRM" to five incumbents with a decade of accumulated third-party validation. My honest advice: don't fight there first.

The smarter move is to find the modifier where the evidence is thin for your whole category and rich for you specifically. Nine times out of ten that's a fit-and-segment modifier. Instead of "best CRM," go for "best CRM for solo real-estate agents." Instead of "best physical therapist," go for "best physical therapist for distance runners recovering from surgery." Narrow the qualified query until you are, plainly, the obvious answer — and then make sure the text says so.

This is the same logic as category creation, just applied at the query level. I would much rather own the definitive answer to "best PT for post-surgery runners" than rank ninth for "best PT." Ranking ninth gets me nothing; AI answers rarely list nine names. Owning the narrow modifier gets me named every time that specific need comes up — and those are the buyers who convert, because the fit is exact. That's real share of the model's attention on a query I can actually win, versus a rounding error on a query I can't.

A worked example: one category, three winners

Take "running shoes." Three buyers, three modifiers, three completely different answers. The first asks for the "best running shoe," and the model returns the consensus darlings — the brands with a decade of ranked lists and reviews behind them. There's almost no dislodging them on the bare superlative, so a newer brand that fights there loses quietly. The second buyer asks for an "affordable running shoe under a hundred dollars," and suddenly half the consensus names fall away because their prices sit too high to be legible as affordable; the winner is whoever has explicit sub-hundred pricing and reviews that use the word "value." The third asks for "the best running shoe for flat feet and overpronation," and now fame barely matters — the model wants the option whose text and reviews specifically connect to that gait problem. A small brand with one authoritative page on overpronation can take that answer outright.

Same noun, three different games. The newer brand that keeps hammering "best running shoe" gets nothing. The same brand that decides to own "best running shoe for overpronators" — with an explicit page, matching reviews, and a mention or two in the right lists — gets named every single time that need surfaces. That's the whole argument in one category: the noun is a crowd, the modifier is a door, and you only need to own the door that's actually yours.

How to build modifier evidence — on your site and off it

On your own site, state the modifier out loud in language a model can lift cleanly. If you want the "affordable" query, publish real pricing — actual numbers, not "contact us for a quote." If you want

"best for small teams," write the sentence "built for teams of ten or fewer" and mean it. Add explicit "who this is for" and "who this is *not* for" sections. The not-for section is quietly powerful: saying "this isn't the right fit for enterprises with complex compliance needs" makes your actual fit unmistakable, and models reward that clarity because it removes ambiguity from the recommendation.

Off your site is where the modifiers you can't self-declare get decided. "Best" and "luxury" live in third-party territory — ranked lists, reviews, editorial coverage, and forum threads. If you want to win those, you have to show up in sources that already use your target adjective. And watch your reviews closely, because they quietly assign modifiers for you: if customers keep describing you as "fast" and "responsive," you'll start winning speed queries whether or not you ever claimed them. The job, then, is to feed the corroboration — earn the mentions, prompt the specific language, and make sure the independent record points where you want it to.

One closing caution, and it's the whole point: you cannot be the cheapest and the most luxurious. You cannot be the fastest and the most meticulous. Every modifier you try to win dilutes the others, because the evidence pulls in opposite directions. Pick the one modifier that matches your actual business, and then make it undeniable in text — on your pages and in the sources that talk about you. The businesses that get named are the ones whose evidence points, without hedging, at a single modifier. That's my read, not a guarantee — but I've yet to see a business win a qualified query while trying to be everything to everyone.

Key takeaways

- ♦ In a recommendation query, the modifier (best, cheapest, luxury, fastest, near me) sets the ranking function — it decides who gets named far more than the category noun does.
- ♦ AI resolves a modifier by mapping it to text-based proxy signals, then favoring sources that make those signals explicit and independently corroborated. It measures how legible your reality is, not reality itself.
- ♦ Almost every buyer modifier falls into five families: superlative, price, prestige, speed/convenience, and fit/segment — and each rewards a different kind of proof.
- ♦ "Best" is the hardest modifier to win because it depends on third-party consensus you don't control; fit-and-segment modifiers are the winnable ones for most businesses.
- ♦ Narrow the qualified query until you are the obvious answer — owning "best CRM for solo agents" beats ranking ninth for "best CRM," because AI answers rarely list nine names.

- ♦ You cannot win opposite modifiers at once; pick the one that matches your real business and make it undeniable on your pages and in the sources that describe you.

Frequently asked questions

Should I try to rank for “best” in my category?

Usually not first. “Best” is decided by independent consensus — ranked lists, reviews, coverage — that you don’t control, so incumbents with years of validation are hard to unseat. Start with a fit-and-segment modifier where you can be the obvious answer, earn corroboration there, and let that authority compound toward the harder superlative queries later.

How does AI decide I’m “affordable” or “luxury” if I never say it?

It reads proxies. For “affordable” it looks for explicit prices and comparative language; for “luxury” it looks at how high your prices sit, your craftsmanship and materials detail, and which other names you’re mentioned alongside. If your text and the sources about you don’t make one of those legible, you won’t be routed to that modifier — so state real numbers and positioning plainly rather than leaving it to inference.

Can one page win multiple modifiers?

Rarely, and often not honestly. Opposite modifiers pull evidence in opposite directions — the cheapest option and the most premium one can’t both be true of the same page, and trying to claim both makes you legible as neither. It’s better to pick the modifier that matches your actual business, own it cleanly, and, if you genuinely serve distinct segments, give each its own focused page rather than one page hedging across all of them.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.