

PLAYBOOK

One Brand, Ten Addresses: Multi-Location AI Visibility

By Scott Tischler · October 7, 2026 · 8 min read

Here is the direct answer: **when a business has several locations, an AI assistant has to decide which one answers a given question — and if the locations are not clearly distinguishable, it will often answer with the brand in general, with the wrong branch, or with a competitor who was easier to resolve.**

This is a different problem from single-location visibility and it is routinely treated as the same one. A single business wants to be recognized. A multi-location business needs to be recognized *and* correctly disambiguated, and the second part is where most of the loss happens.

The three failures

The locations are interchangeable. Every location page is the same page with the city swapped. A model cannot tell them apart because there is nothing to tell apart. Asked about the one in a particular suburb, it has no more to say than it does about any other, so it answers generically.

The brand swamps the branches. All the authority accrues to the main domain and none to any specific location. The brand is recognized; no individual location is. Fine for brand questions, useless for the question that actually converts, which is almost always local.

The branches compete with each other. Two locations both trying to answer the same regional question, splitting the corroboration between them, and neither accumulating enough to win it. This is the quiet one, because nothing looks broken.

What actually distinguishes a location

Not the city name in the heading. That is the minimum and it does almost nothing on its own.

What distinguishes a location is the stuff that is only true there: the staff who work at that site, named; the hours that location actually keeps, which usually differ; the services or equipment available there

and not elsewhere; the parking situation, the access, the room; the neighborhoods it genuinely serves; the reviews written by people who went to that address.

Put plainly: a location page should contain things that would be false if you pasted them onto another location's page. If everything on the page would survive that paste, the page is not describing a location. It is describing a brand with a city in the title.

Reviews are the strongest differentiator and the most ignored

Reviews are tied to a specific address and written by people who were physically there. They mention the staff by name and the particulars of that site. Nothing you write about your own location carries the same weight as a stranger describing it.

Which means review distribution across locations matters, not just review volume overall. A brand with four hundred reviews concentrated at one branch and eleven at another does not have a strong presence at the second one, whatever the brand-level average says. Worth checking per location rather than in aggregate, because the aggregate hides exactly the problem you need to see.

Structure that helps a machine

- A distinct, indexable page per location, each with its own address, hours, phone number and named staff.
- Location markup on each, with the address and opening hours stated as data rather than inferred from prose.
- A locations index that links to every one of them, so the set is discoverable as a set.
- Consistent naming across every platform. "Brand — Lehi" everywhere beats "Brand Lehi", "Brand of Lehi" and "Brand (Lehi)" scattered across directories, because inconsistent naming is how one entity becomes three half-entities.
- A clear statement, per location, of the area it serves — which also resolves the branches-competing problem by assigning the regional question an owner.

Deciding which location owns which question

This is a business decision before it is a technical one, and it is the one most teams never make explicitly.

For any question that spans your whole region, pick the location that should own it and let the others point there. Splitting it between two is how you lose it to a competitor who only has one address and

therefore never had to decide. A single clear owner concentrates the corroboration instead of dividing it.

Where the leverage is

If I had limited time with a multi-location business, I would spend it making each location page contain things that are true only at that location, and checking the per-location review distribution rather than the brand average. Those two together address all three failure modes, and both are achievable without a rebuild.

The deeper point: with multiple locations, part of your competition is yourself. Every ambiguity between your own branches is effort spent confusing a system that was trying to recommend you.

Key takeaways

- ◆ Multi-location businesses must be recognized and correctly disambiguated; the second part is where most loss occurs.
- ◆ If a location page would still be true pasted onto another location, it is not describing a location.
- ◆ Name the staff, state the real hours, and list what exists only at that site.
- ◆ Check review distribution per location, not the brand average, which hides weak branches.
- ◆ Keep naming identical across every platform; inconsistent naming splits one entity into several weak ones.
- ◆ Assign each regional question a single owning location so your branches stop splitting corroboration.

Frequently asked questions

Why does AI sometimes answer with my brand instead of a specific location?

Usually because no individual location has accumulated enough distinguishing information to be resolved on its own. The authority sits on the main domain and the location pages are interchangeable, so the brand is the only thing the system can confidently name.

Should each location have its own page or is one locations page enough?

Each location needs its own indexable page with its own address, hours, phone number, named staff and location markup, plus an index page linking the full set. A single combined page gives a model nothing to tell the locations apart with.

How do I stop two of my locations competing for the same search?

Decide which one owns that regional question and have the others point to it. Splitting the question between two locations divides the corroboration so neither wins it, and a single-location competitor who never had to make that choice takes it instead.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.