

AI & SEARCH

Multimodal Search: Getting Recommended When Customers Search With a Photo

By Scott Tischler · September 18, 2026 · 8 min read

A growing number of searches don't start with words anymore. Someone photographs a pair of shoes on the street, screenshots a product from a video, or points their phone at a broken part under the sink and asks the AI “where can I get this” or “who can fix this.” The interface changed and most businesses haven't noticed, because they're still optimizing for typed phrases while a real slice of demand now arrives as pixels. If you sell anything visual — products, spaces, food, physical work — this is worth understanding now rather than after it's obvious.

Here's the direct answer to how it works: an AI handling an image search runs two steps, not one. First it **identifies** what's in the image and turns it back into words. Then it **recommends** using those words, exactly the way it would with a typed query. Almost every business that loses in multimodal search loses at the first step, and the fix for the first step is different from anything in a normal SEO or AEO playbook.

What actually happens when someone searches with an image

When you hand an AI a photo, it can't recommend a picture. It has to convert what it sees into language it can reason about — “a white ceramic pendant light with a brass stem, mid-century style” or “a cracked P-trap under a bathroom sink” or “a storefront that looks like a specialty coffee roaster.” That conversion is the identification step, and everything downstream depends on it being right. Once the model has words, it's back on familiar ground: it runs essentially the same assembly process it uses for any recommendation, pulling names it can corroborate and describe for that now-textual query.

So the uncomfortable truth is that multimodal search still gets *decided* in text. The image just determines *which* text query you end up competing in. Get identified as the wrong thing and you're competing for a customer who was never going to buy from you. Get identified precisely and you're in a query you can actually win with the authority work you'd do anyway.

The identification step is where most businesses lose

The model identifies an image using whatever visual and contextual evidence it can find. For a product, that means the model does better when your item is visually documented on the open web in a way that ties the picture to a name: clear product photos from multiple angles, associated with a specific product name and category, ideally echoed across more than one source. If the only images of your product are on a single page with a vague caption, the model has thin evidence to match against and will often round you off to a generic category — or to a bigger competitor whose product it *can* identify confidently.

For a physical business — a restaurant, a studio, a storefront — identification leans on the association between images of your place and your name and location across the web: your listings, your reviews with photos, mentions that pair a picture with who and where you are. A business that's photographed often and consistently labeled is easy to identify. A business with two stock photos and no visual footprint is a coin flip.

There's a compounding effect worth naming. Every correctly-labeled image of your product or place doesn't just help with that one photo — it strengthens the model's overall association between your look and your name, which makes the *next* identification easier too. A business with a rich, consistent visual footprint builds a kind of recognition equity, where even a partial or awkward customer photo still resolves to it. A business with a thin footprint stays a coin flip indefinitely, because there's never enough evidence for the association to firm up. This is why I treat visual documentation as an asset you accumulate, not a task you check off.

I want to be honest about the limits of my confidence here: the exact weighting of these systems isn't public and it's moving fast. But the direction is not subtle. The businesses that are easy to identify from an image are the ones whose visual presence on the web is *abundant, labeled, and consistent*. That's the lever.

Text still decides the recommendation

Once you've been correctly identified, the second step is the familiar one, and it rewards the same fundamentals I write about constantly. The model recommends the identified thing based on corroboration, describability, recency, and fit to any qualifier in the ask. “Where can I buy this pendant light” becomes a sourcing question, and it'll favor sellers it can confirm actually carry it. “Who can fix this” becomes a local-service question, and it'll favor providers whose specialty and service area are clearly stated and independently confirmed.

This is actually good news. It means you don't need a separate strategy for the recommendation half of multimodal search — the authority, review, and consistency work that makes you recommendable in

text already carries over. What's genuinely new, and what most people are neglecting, is the identification half.

The playbook for visual visibility

If image-started searches matter for your category, here's where I'd put the effort.

- **Document your products and spaces abundantly.** Multiple clear images, multiple angles, good lighting, on more than one property — your site, your listings, review platforms, wherever your things legitimately appear. Sparse visual evidence is the single most common reason a business can't be identified.
- **Label every image in plain language.** Tie each picture to a specific name, category, and the qualifiers that matter, in captions, alt text, and surrounding copy. You're teaching the model the words to convert the image into.
- **Keep visual details consistent with your text.** If your photos show one thing and your description says another, identification gets shaky. The picture and the words about it should agree.
- **Encourage customer photos.** Reviews and posts with images of your product or place are powerful corroboration — independent visual evidence that pairs your look with your name.
- **Then do the ordinary authority work.** Once identification is solid, everything that makes you recommendable in text — corroborated specialty, clear service area, freshness, consistency — does the rest.

How this plays out by what you sell

The identification problem looks different depending on your category, and it's worth being specific about your own. If you sell **distinctive physical products** — furniture, jewelry, apparel, specialty equipment — you have the most to gain and the most to lose, because a well-documented item can be matched to your name from a single photo, while an under-documented one gets handed to whichever competitor the model *can* confidently name. Abundant, labeled product imagery is close to non-negotiable here.

If you run a **local physical business** — a restaurant, gym, clinic, or shop — identification rides on the association between images of your place and your name and location. The work is making sure your storefront, interior, and offerings are photographed and correctly labeled across your listings and reviews, so a snapshot of your sign or your space resolves to you rather than to a generic “coffee shop near here.” If you provide a **repair or trade service**, the customer's image is usually of the *problem*, not of you — a leaking valve, a cracked screen, a dead furnace. There, the identification step is about the object, and your job is on the recommendation side: make it unmistakable that you fix exactly that

thing in exactly that area. And if you sell something **purely abstract** — software, consulting, finance — image search barely touches you today, and you can spend your attention elsewhere with a clear conscience.

My honest read on where this goes

I don't think image-first search replaces typed search; I think it becomes a normal on-ramp alongside it, especially on phones and especially for anything you can see before you can name. That makes visual documentation a form of visibility infrastructure, the same way a clean, consistent web presence became infrastructure for text. The businesses that get ahead won't be the ones with the most artistic photography — they'll be the ones whose products and places are documented so thoroughly and labeled so plainly that a machine can look at a customer's blurry phone snapshot and know, without guessing, exactly who makes that and where to get it. That's an unglamorous, do-it-once kind of project, which is exactly why the businesses that do it will quietly own the queries nobody else prepared for.

Key takeaways

- ♦ A real and growing share of searches now start with an image or screenshot, and most businesses are still optimizing only for typed phrases.
- ♦ AI handles an image search in two steps: it identifies what's in the picture and turns it into words, then recommends using those words like any text query.
- ♦ Most businesses lose at the identification step — thin or unlabeled visual evidence gets them rounded off to a generic category or a bigger competitor.
- ♦ Identification improves when your products or spaces are documented abundantly, labeled in plain language, and shown consistently across multiple sources.
- ♦ The recommendation step is the familiar one — corroboration, describability, recency and fit — so your existing authority work carries straight over.
- ♦ Treat visual documentation as visibility infrastructure: do it thoroughly once and you own image-started queries competitors never prepared for.

Frequently asked questions

Does image search really matter for my business, or is it a novelty?

It depends on whether your category is something people encounter visually before they can name it — products, food, spaces, physical repairs, styles. For those, image-started search is already a normal on-ramp, especially on phones. For purely abstract services it matters less. My advice is to judge it by your category, not by the hype, but if you sell anything you can photograph, it's worth preparing for now.

What's the single most important thing for showing up in multimodal search?

Making yourself easy to identify from an image. That means abundant, clearly labeled visual documentation of your products or space across more than one source, so the model has strong evidence to match a customer's photo against your name. Get identified correctly and the rest of the process is the same authority work that already makes you recommendable in text.

Is alt text still worth writing if AI can “see” the image directly?

Yes. Even when a model can interpret an image, the words around that image — alt text, captions, nearby copy — are what tie the picture to a specific name, category, and set of qualifiers. That labeling is exactly the language the model converts the image into, so it's not redundant with visual recognition; it's what makes the recognition land on *you* instead of a generic description.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.