

PLAYBOOK

# Original Data Is the AEO Asset Almost Nobody Publishes

By Scott Tischler · September 20, 2026 · 8 min read

---

**If you want AI to cite you, publish a number nobody else has. Answer engines are built to attach a source to a specific claim, and original data — a stat, a benchmark, a survey result that only you own — is the most citable thing you can put on the internet. It's also the one asset almost nobody in your category bothers to make.**

I spend my days watching how ChatGPT, Gemini, Perplexity and Google's AI answers decide who to name. The pattern that keeps repeating is simple: when the model states a fact that needs backing, it reaches for the page that *is* the backing. Opinion pages get paraphrased and forgotten. Pages that contain the actual figure — "the average X is Y," "43% of Z do this" — get cited by name, because the model needs somewhere to point when a person asks "says who?"

Most businesses are competing to write the best explanation of a topic everyone already agrees on. That's a losing game, because there are a thousand good explanations and the model can synthesize its own. Original data is different. There's exactly one source for a number you generated, and if the number is useful, the model has no choice but to send people to you.

## Why answer engines are biased toward data

Under the hood, most AI answers are built by retrieving a handful of sources and then generating a response grounded in them. The generation step is looking for things it can assert with confidence. A vague claim ("many companies struggle with this") can be written from the model's own training. A precise claim ("the median onboarding time is 11 days") cannot — it has to come from a document. That's the moment a citation gets created.

So the mechanism rewards specificity. The more a sentence sounds like it came from a measurement rather than an opinion, the more likely the model is to treat the underlying page as a source worth naming. This is why a mediocre survey with real numbers often out-cites a brilliant essay with none. My

read, not a guarantee: the models are pattern-matching for "this looks like evidence," and a number is the strongest evidence signal a page can send.

There's a second effect. Once one AI system cites your data, that citation becomes part of the corpus other systems learn from. A figure that gets picked up, quoted, and re-quoted starts to look like a settled fact — and you're the settled fact's origin. That compounding is the real prize. You're not just getting cited once; you're becoming the entity the number is attached to.

## What counts as "original data" — and what doesn't

You don't need a research department. Original data just means a number you produced that didn't exist in that form before. A few kinds that are within reach of almost any operator:

- **Your own operational numbers, anonymized.** If you run a business, you're sitting on data. Average project timeline, typical cost range, before-and-after results across your clients, the percentage of customers who ask for a specific thing. Aggregate it, strip anything identifying, and publish the pattern.
- **A small survey of your audience or industry.** Two hundred honest responses to five sharp questions beats zero responses to a perfect study. Ask something people actually wonder about and report what you find.
- **A benchmark or teardown.** Test a set of tools, products, or approaches under the same conditions and publish the comparison table. The methodology being consistent matters more than it being fancy.
- **A tracked trend over time.** Measure the same thing every quarter and publish the movement. "Here's how this changed over 18 months" is a story only the person who kept measuring can tell.

What doesn't count: repackaging someone else's statistic, quoting a study everyone already quotes, or dressing up an opinion as a finding. If a person could get the same number from ten other pages, the model can too, and you've added nothing worth pointing at.

## How to publish data so AI can actually use it

Generating the number is half the job. The other half is presenting it so an answer engine can extract it cleanly and trust it. Here's how I'd structure it.

**Put the headline number in a single, self-contained sentence.** "In our survey of 312 franchise owners, 61% said they'd chosen their brand before comparing more than two options." That sentence can be lifted out and dropped into an answer without losing its meaning — which is exactly what you want to happen.

**Show your methodology in plain sight.** Say how many people, over what period, measured how. Answer engines are increasingly weighing whether a source looks credible, and a visible method is a credibility signal. It also protects you: a number with a stated method is one you can defend when someone asks where it came from.

**Give the data a stable home and a real title.** A page titled "2026 [Your Industry] Benchmark Report" that stays at one URL becomes a fixed reference the model can return to. Numbers buried three paragraphs into an unrelated blog post get missed.

**Make the raw finding easy to quote and hard to misread.** Round sensibly, state the unit, and don't hedge the number itself into vagueness. You can be honest about limitations in the next sentence — "small sample, directional only" is fine and even helps — but the figure should be clean.

**Refresh it on a schedule and date it.** Data decays. A benchmark labeled 2024 loses pull as the year turns. If you commit to updating it, you own the citation year after year instead of handing it to whoever measures next.

## **A worked example: turning what you already track into a stat**

Let me make this concrete, because "publish original data" can sound like it requires a budget you don't have. Say you run a service business — a studio, an agency, a contractor, a clinic. You already track things: how long jobs take, what people ask for, why they say no, what they end up spending. That's raw material sitting in your calendar, your invoices, and your inbox.

Pick one question your prospects genuinely have. "How long does this actually take?" is a good one, because every buyer wonders and most sellers answer vaguely. Pull your last hundred completed jobs, calculate the real median and range, and write it up: "Across our last 100 projects, the typical timeline was 18 days, with 80% landing between 12 and 28 days." Add a sentence on what pushes a project to the long end. You now own a number that answers a real question, and you got it from records you already keep.

Do the same with the objections you hear, the price ranges you quote, or the outcomes you deliver. None of this requires new measurement — it requires paying attention to measurement you're already doing and being willing to publish the honest result. The businesses that win here aren't the ones with the most data. They're the ones willing to count what they already have and put the number where an AI can find it.

## **The honest caveats**

This isn't a cheat code, and I won't pretend it is. A single stat page won't outrank an established authority overnight, and if your number is boring or nobody's asking the question it answers, it'll sit there uncited. The data has to be genuinely useful to a real question people ask an AI. You also have to be willing to be accurate — publishing a shaky number that gets picked up and then debunked does more damage than staying quiet.

And there's a discipline cost. The businesses that win here are the ones that treat data as a product they maintain, not a one-time stunt. That's actually good news, because it means the barrier is effort, not budget — which is the kind of barrier a determined operator can clear when their competitors won't.

The category-level takeaway: everyone is writing content to explain things AI can already explain. Almost nobody is manufacturing the facts AI needs and can't generate on its own. Be the source of a number, and you stop competing for attention — you become infrastructure the answers are built on.

### Key takeaways

- ♦ Answer engines create citations at the moment they state a specific fact, so a page that owns a number is far more citable than one that only explains a concept.
- ♦ You don't need a research department: anonymized operational numbers, a small audience survey, a consistent benchmark, or a tracked trend all count as original data.
- ♦ Repackaged or widely-available statistics add nothing — if the model can get the same number from ten other pages, it won't need yours.
- ♦ Present data for extraction: a self-contained headline sentence, a visible methodology, a stable dated URL, and a clean, quotable figure.
- ♦ Refresh and re-date your data on a schedule so you keep owning the citation year after year instead of losing it to whoever measures next.
- ♦ The winners treat data as a maintained product, not a one-time stunt — the barrier is effort, not budget, which is why so few competitors clear it.

## Frequently asked questions

### Do I need a large study for AI to cite my data?

No. A small, honest dataset with a clear method often out-cites a large vague one. Two hundred real survey responses, or your own anonymized operational averages, can be enough — what matters is that the number is specific, useful to a real question, and clearly your own.

## Why does original data get cited more than a well-written explanation?

Because AI can generate its own explanation of a known topic, but it cannot invent a specific number — it has to attribute that to a source. When an answer states a precise fact, it reaches for the page that contains that fact, which is how the citation gets created. Explanations get paraphrased; data gets attributed.

## How should I format a data page so answer engines can use it?

Put the headline finding in one self-contained sentence, show your methodology in plain sight, give the page a real title and a stable URL, keep the figure clean and quotable, and date it. Then update it on a schedule so the number stays current and you keep owning the citation.

**About the author.** Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.