

HOW AI WORKS

Query Fan-Out: How One AI Question Becomes a Dozen Hidden Searches

By Scott Tischler · September 30, 2026 · 9 min read

When you type one question into an AI search tool, it almost never runs just that one search. Behind the scenes it quietly rewrites your question into a whole set of related searches, runs them all, and then stitches the results together into a single answer. That behavior has a name — **query fan-out** — and if you want to understand why your business shows up for some AI answers and vanishes from others, this is the mechanism to understand first.

I spend most of my week watching how AI assistants decide who to cite and recommend, and query fan-out is the piece most people miss. They optimize for the exact phrase a customer might type, get frustrated when they don't appear, and never realize the model wasn't searching for that phrase at all. It was searching for eight or ten other things it invented on the spot. This is my working model from observing thousands of answers, not a leak of any lab's internal architecture — but it's consistent enough that I plan around it.

What query fan-out actually does

Traditional search took your keywords more or less literally and returned a ranked list of pages. You did the work of reading, comparing, and deciding. AI search inverts that. Instead of handing you ten links, the system tries to hand you one synthesized answer — and to build that answer confidently, it needs to look at the question from several angles at once.

So it decomposes. A single prompt like "what's the best CRM for a small law firm" gets silently expanded into a fan of sub-questions: best CRM for small businesses, CRM features law firms need, CRM pricing for solo attorneys, alternatives to the obvious market leader, reviews of legal-specific CRMs, and so on. Each of those becomes its own retrieval. The model gathers documents across all of them, weighs what it found, and composes an answer that draws from the strongest, most corroborated material across the entire fan.

The practical consequence is enormous. You are no longer competing to rank for *one* query. You're competing to be a strong, retrievable answer across a whole neighborhood of related queries you never see — and any one of them can be the door you get pulled in through, or the door you get shut out of.

Why this breaks traditional keyword thinking

The old SEO instinct is to find the phrase with volume and optimize a page to match it. Query fan-out makes that instinct actively misleading, for a few reasons.

- **You can't see the queries that matter.** The sub-questions the model generates aren't in any keyword tool. They're synthesized in the moment from the user's intent, the conversation so far, and the model's own sense of what a complete answer requires. Optimizing only for the visible phrase means optimizing for one branch of a tree you can't map.
- **Coverage beats exact-match.** A page that answers one query perfectly can lose to a source that answers the surrounding cluster decently. If the model fans out into ten sub-questions and you're a strong match for one while a competitor is a solid match for seven, the competitor gets synthesized into the answer and you don't.
- **The answer is a composite, not a winner.** There's rarely a single "position one." The response is assembled from many retrievals, so the goal shifts from "rank first for the keyword" to "be reliably present across the fan." Presence across the cluster is the new ranking.

I've watched businesses with a beautifully optimized flagship page get skipped entirely because they'd bet everything on one query and had nothing to say about the adjacent questions the model actually asked. The page was excellent. It just wasn't *enough surface area*.

How to figure out your fan

You can't see the model's internal sub-queries directly, but you can reconstruct the shape of the fan well enough to work with. Here's the approach I use.

Start from the intent, not the phrase. Take the real question a customer is trying to answer and ask yourself what a thorough human researcher would look up to answer it completely. Those are your sub-questions. For "best CRM for a small law firm," that's features, pricing, alternatives, integrations, reviews, migration, and legal-specific needs. That list is a rough map of the fan.

Ask the AI tools to expand it for you. Prompt ChatGPT, Gemini, and Perplexity with your core question and then ask, "what related questions did you consider to answer that?" or "what would

someone need to know to choose well here?" You'll get a serviceable approximation of the decomposition. It's not the literal internal fan, but it's close enough to reveal the clusters you're missing.

Watch the follow-ups. The questions users naturally ask next — the second and third turns in a conversation — are often the same territory the fan already covered. If people always ask "how much does that cost?" right after the recommendation, cost is in the fan. Mine your own sales conversations and support tickets for these.

Look at what the answers cite. When an AI answers your core question, note which sources it pulls from and what sub-topics they cover. The shape of what got cited tells you the shape of what got searched.

How to actually win across the fan

Once you accept that you're competing across a cluster rather than a keyword, the content strategy changes. The goal is to be a strong, retrievable, corroborated answer to as many branches of the fan as possible.

Build topical depth, not just a landing page. Instead of one page straining to rank for the money keyword, cover the whole neighborhood — the features question, the pricing question, the comparison question, the "who is this right for" question. Each becomes a retrievable answer to a different branch. Depth across the cluster is what gives the model multiple places to find you.

Answer each sub-question cleanly and self-containedly. Fan-out retrieval works at the passage level, not the whole-page level. A section that directly and completely answers one sub-question — with the answer stated up front, in plain language, not buried under three paragraphs of throat-clearing — is far more retrievable than the same information scattered across a page. Write sections that can stand alone, because that's how they'll be pulled.

Cover the comparison and alternatives branches on purpose. The fan almost always includes "alternatives to X" and "X vs Y" sub-questions, and most businesses refuse to touch them because it feels like promoting competitors. That refusal is why they're absent from exactly the branch where buying decisions get made. Address comparisons honestly and you show up in the part of the fan that converts.

Get corroborated across the cluster. Being the answer to a sub-question on your own site is a claim. Being described consistently across independent sources — reviews, directories, earned media, third-party roundups — is what makes the model confident enough to synthesize you into the answer. Corroboration doesn't just help one query; it helps every branch of the fan where your name comes up.

A concrete example

Say you run a boutique CRM built specifically for law firms. A prospect asks an AI, "what's the best CRM for a small law firm?" You've got a gorgeous homepage optimized for exactly that phrase. And you lose anyway.

Here's what happened inside the fan. The model expanded the question into sub-searches for legal-specific features, pricing for small practices, alternatives to the big generic players, integration with legal billing software, and reviews from actual attorneys. Your homepage matched the headline query well — but you had no clear, standalone content on legal billing integrations, no honest comparison against the generic incumbents, and no independent reviews the model could corroborate. A competitor did. So the competitor got assembled into the answer across five branches while you matched one, and the synthesized recommendation went to them.

The fix wasn't a better homepage. It was building genuine, retrievable answers to the branches you were losing — the integration page, the honest comparison, the case for the specific kind of firm you serve — and earning independent mentions that corroborated all of it. Spread across the fan instead of stacked on one query, the same business becomes hard to leave out.

The mindset shift

Query fan-out is the clearest reason the old "rank for the keyword" model is ending. You're not trying to win a single search anymore. You're trying to be a dependable, well-corroborated presence across a whole invisible cluster of related searches — any of which can be the reason you make it into the answer.

My honest read is that fan-out is only going to get wider and deeper as these systems get better at decomposing hard questions. The businesses that win won't be the ones with the single most optimized page. They'll be the ones with real depth across the cluster and enough independent corroboration that the model can confidently pull them in through whichever door the user happened to open. Stop optimizing for the query you can see. Start earning your place across the fan you can't.

Key takeaways

- ♦ AI search rarely runs your literal query — it fans it out into many synthesized sub-searches (query fan-out) and stitches the results into one answer.
- ♦ You're no longer competing to rank for one keyword; you're competing to be a strong, retrievable answer across a cluster of related questions you never see.

- ♦ Coverage beats exact-match: a source that decently answers seven branches of the fan beats one that perfectly answers a single query.
- ♦ Reconstruct the fan by starting from intent, asking the AI tools what related questions they considered, mining follow-ups, and studying what answers actually cite.
- ♦ Win the fan with topical depth, clean self-contained sections per sub-question, honest comparison/alternatives content, and independent corroboration.
- ♦ Retrieval happens at the passage level — sections that answer one sub-question completely and up front are far more likely to be pulled than the same facts scattered across a page.

Frequently asked questions

What is query fan-out in AI search?

Query fan-out is when an AI system takes your single question, rewrites it into a set of related sub-questions, runs a search for each of them, and then synthesizes the results into one answer. So a prompt like 'best CRM for a small law firm' quietly becomes searches for features, pricing, alternatives, integrations and reviews — and the final answer is assembled from the strongest material found across all of them, not from one literal search for your phrase.

How do I optimize for queries I can't even see?

You reconstruct the shape of the fan instead of guessing exact phrases. Start from the real intent behind the question and list what a thorough researcher would need to look up. Ask ChatGPT, Gemini and Perplexity what related questions they considered when answering. Mine the follow-up questions your customers actually ask. Then build clean, self-contained answers to each of those sub-questions and earn independent sources that corroborate them, so you're retrievable across the whole cluster rather than betting on one keyword.

Does query fan-out make traditional SEO useless?

Not useless, but insufficient on its own. The foundational work — crawlable pages, clear structure, genuine authority — still matters because it's what makes your content retrievable in the first place. What changes is the target: instead of optimizing one page for one high-volume keyword, you build depth across a cluster of related questions and get corroborated across independent sources. Think of it as SEO's discipline applied to a much wider surface area than a single query.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.

