

STRATEGY

The Sameness Problem: Why AI Can't Tell You Apart From Your Competitors

By Scott Tischler · September 29, 2026 · 9 min read

If you and five competitors all say the same things — "trusted," "experienced," "customer-focused," "results-driven" — an AI has no reason to pick you over any of them. When the model can't tell you apart, it falls back to whatever it *can* measure: size, popularity, and how many other sources happen to mention you. The fix isn't better adjectives. It's a wedge — one specific, verifiable thing you're the obvious answer for that your competitors can't claim.

This is my read as a practitioner, not a guarantee about how any single model works under the hood. But I've watched it play out enough times across ChatGPT, Gemini, Perplexity and Google's AI answers that I now treat "sameness" as one of the most common and least-discussed reasons good businesses stay invisible. They're not being penalized. They're being averaged.

Sameness is a visibility problem, not just a branding one

Marketers have talked about differentiation for decades, but they've mostly framed it as a *persuasion* problem: stand out so the human buyer remembers you. AI changes the stakes, because now something has to distinguish you before a human ever sees the shortlist. The model builds the consideration set first. If you don't survive that step, your positioning never gets a chance to work.

Here's the mechanism as I understand it. When someone asks an AI "who's the best X in Y," the model isn't pulling a ranked list off a shelf. It's assembling an answer from patterns in what it retrieved and what it learned in training. To name a specific business, it needs a reason — some attribute, association, or corroborated claim that connects *you* to *that question* more strongly than it connects your competitors. If everyone in your category is described in identical language, that connective tissue doesn't exist for anyone in particular, and the model does the safe thing: it reaches for the names that are simply bigger or more talked about.

So sameness doesn't just make you forgettable to buyers. It makes you mathematically hard to select. You've given the model nothing to grip.

What AI does when everyone looks the same

When a category is undifferentiated, the model still has to answer the question, so it falls back on the signals that are always available. In my experience those are, roughly in order:

- **Raw prominence.** Who gets mentioned the most across the web, regardless of why. Bigger brands, older brands, and heavily-covered brands win by default because volume is the easiest signal to read.
- **Corroboration.** Who shows up on the third-party "best of" lists, directories, and roundups. If ten sources independently mention a competitor for your exact category, that's a stronger reason to name them than anything you say about yourself.
- **Recency and activity.** Who looks alive right now — publishing, being cited, getting reviewed. A stale but once-prominent name can hold on for a while, but freshness breaks ties.
- **Category fit.** Whose described focus most tightly matches the specific query. This is the one you can actually engineer, and it's where the wedge lives.

Notice that three of those four favor whoever is already big. If you're the smaller or newer player and you compete on the same generic ground as everyone else, you are playing the exact game the incumbent wins. Sameness hands the answer to the market leader.

The wedge: own a narrower question

The way out is not to be better at the broad question. It's to be the unambiguous answer to a narrower one. I call this the wedge, and the whole move is to trade a big query you'll lose for a smaller query you can own outright.

"Best marketing agency" is a fight you can't win against a thousand competitors and every national brand. "Marketing agency for independent dental practices in the Southeast" is a question with maybe a handful of legitimate answers — and if you're one of them and you've made that focus legible, you can be the name the model reaches for every time. You didn't get more visible in general. You got completely visible in a specific place.

A good wedge has three properties. It's **specific** enough that only a few businesses could honestly claim it. It's **true** — you actually deliver on it, so the corroborating evidence exists or can be built. And it maps to a **question real people ask**, because a distinction nobody searches for is just trivia. The art is finding the intersection of "what I'm genuinely best at" and "what someone would actually type into an AI." When those overlap, you've found your wedge.

The counterintuitive part is that narrowing usually *grows* your visibility rather than shrinking it. People worry that picking a niche means turning away business, but in AI answers the opposite tends to happen. Owning "personal-injury intake automation for solo firms" doesn't lock you out of the broader "legal tech" conversation — it earns you a foothold in it, because the model now has a concrete, corroborated reason to associate you with the category at all. Specific authority is how you back into general authority. You start as the definitive answer to a small question and let that pull you into the bigger ones. Trying to run it in reverse — claim the big category first and hope specificity follows — is exactly the sameness trap that got you averaged in the first place.

Make the difference legible to a machine

Finding the wedge is half of it. The other half is making it retrievable, because a distinction that lives only in your founder's head or in a slide deck doesn't help the model at all. The rule I use: if you can't point to the specific, corroborated evidence of a claim, assume the AI can't either.

That means saying the specific thing plainly and often — on your own site, in the language a customer would use, not in marketing abstraction. It means the claim showing up in *other people's* content: the roundups, the interviews, the case studies, the directories, the podcast where you explained exactly who you serve. Sameness is broken by corroboration, and corroboration you don't control is worth more than anything on your homepage.

Concreteness is what makes it stick. "We serve dental practices" is soft. "We've run patient-acquisition campaigns for 40+ independent dental practices across Georgia, Florida and the Carolinas" is a claim a model can attach to a query, because it's specific, it's countable, and it's the kind of thing third parties will repeat almost verbatim. Numbers, named niches, and concrete proof are the raw material of differentiation a machine can actually use.

The traps: fake differentiation and adjective inflation

Two failure modes show up constantly when people try to fix this, and both make it worse.

The first is **adjective inflation** — trying to stand out by turning up the volume on the same generic claims. "The most trusted." "Uniquely positioned." "Truly customer-obsessed." These add words without adding signal, because every competitor is saying them too and none of them are verifiable. To a model, they're noise. If a claim can't be checked and everyone makes it, it does nothing to distinguish you.

The second is **fake differentiation** — inventing a distinction that's technically unique but that nobody's actually looking for, or that you can't back up. You can absolutely become "the only firm using our proprietary 7-Pillar Clarity Framework™," and the model may even repeat it, but if no one searches for

that and no third party corroborates it, you've won a category of one that generates zero recommendations. A wedge only works when it sits on top of real demand and real evidence. Manufacture either one and you've built a distinction that doesn't convert.

How to tell whether your wedge is working

You don't have to guess. Build a short prompt set — the actual questions a customer would ask an AI in your category, including the narrower ones your wedge is meant to own — and run it across the major engines on a regular cadence, logging what comes back. Watch two things: whether you appear at all, and *who appears alongside you*. If you're still landing in the same undifferentiated cluster as everyone else on the broad queries, that's expected — but on your wedge queries, you should be moving toward being named early and consistently.

When the wedge is working, you'll see it first on the narrow questions, then watch it slowly earn you a mention on the broader ones too, because the model has learned a specific reason you belong in the conversation. That's the whole strategy in one sentence: stop trying to be a slightly better version of everyone else, and become the only honest answer to a question worth asking. Sameness is the default. A wedge is a decision.

Key takeaways

- ◆ Sameness isn't just a branding problem — it's a selection problem. If you and your competitors are described in identical language, the model has nothing to grip and can't single you out.
- ◆ When a category looks the same, AI falls back on prominence, corroboration, and recency — signals that structurally favor whoever is already biggest. Sameness hands the answer to the market leader.
- ◆ The way out is a wedge: stop competing on the broad query you'll lose and become the unambiguous answer to a narrower one you can own outright.
- ◆ A real wedge is specific enough that few can honestly claim it, true enough that evidence exists, and mapped to a question real people actually ask an AI.
- ◆ A distinction only counts if a machine can retrieve it — say it plainly on your own site, and get it corroborated in other people's content, in concrete, countable terms.
- ◆ Avoid adjective inflation (louder generic claims) and fake differentiation (a unique-but-irrelevant angle). Both add words without adding a signal the model can use.

Frequently asked questions

Isn't differentiation just a branding exercise?

It used to be mostly about persuasion — stand out so a buyer remembers you. With AI in the loop, something has to distinguish you before a human ever sees the shortlist, because the model builds the consideration set first. If you don't survive that step, your positioning never gets its chance. So differentiation is now also a retrieval and selection problem, not only a branding one.

What makes a good wedge versus a gimmick?

A good wedge has three properties at once: it's specific enough that only a few businesses could honestly claim it, it's genuinely true so corroborating evidence exists or can be built, and it maps to a question real people actually ask. A gimmick usually fails the last two — it's a unique-sounding label nobody searches for and no third party repeats, so it generates zero recommendations.

How do I know if my wedge is actually working in AI answers?

Build a small prompt set of the real questions customers ask in your category, including the narrower ones your wedge targets, and run it across ChatGPT, Gemini, Perplexity and Google's AI answers on a regular cadence. Log whether you appear and who appears alongside you. You should see movement on the wedge queries first — named earlier and more consistently — before it earns you mentions on the broader ones.

About the author. Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.