

HOW AI WORKS

# What Does AI Think You're Known For? The Attribute-Extraction Problem

By Scott Tischler · September 24, 2026 · 8 min read

---

**When an AI recommends a business, it doesn't just decide whether to name you — it decides what to say *about* you. It attaches attributes: fast, affordable, premium, reliable, great for beginners, best for enterprise. That short list of adjectives is what I call the attribute layer, and it often matters more than the citation itself, because it's what actually routes a buyer toward or away from you. This is my read as a practitioner, not a claim about any model's internals, but the pattern is consistent: AI decides what you're "known for," and it decides it from your footprint, not your intentions.**

Most businesses obsess over whether they get mentioned. Fewer ask the more important question: when I do get mentioned, what am I mentioned as? You can be named in an answer and still lose, because the two words the model attached to your name sent the reader to your competitor.

## The difference between being named and being characterized

Think about how you actually use an AI to make a decision. You don't just ask "who does X." You ask "who's the best for someone like me," or "what's the affordable option," or "which one is easiest to get started with." The model answers by naming businesses *and* characterizing them — "A is the premium pick, B is the budget-friendly one, C is best if you want hands-on support."

Those characterizations are the attribute layer. They're the difference between showing up and showing up as the right thing. A jeweler named as "a solid mid-range option" and a jeweler named as "the go-to for custom, one-of-a-kind pieces" have completely different outcomes even if they appear in the exact same answer. Attribute extraction is the process by which the model decides which of those you are.

And here's the uncomfortable part: you don't get to declare your attributes. The model infers them. It reads the web about you — your own pages, yes, but also reviews, third-party descriptions,

comparisons, forums — and it distills a consensus of what you're associated with. If that consensus is thin, generic or wrong, that's what you're "known for," regardless of how you'd describe yourself.

## Where AI gets your attributes from

In my experience, the attributes a model attaches to you come from a few layers of signal, and they don't all carry equal weight.

- **Repeated language across independent sources.** This is the heavy hitter. When many different people and pages describe you with the same words — "fast turnaround," "patient with beginners," "pricey but worth it" — that repetition reads as consensus, and consensus is what the model trusts. One page saying you're the best is a claim; twenty independent sources implying it is a fact, as far as the model is concerned.
- **The words your customers actually use.** Reviews and forum posts are gold here, because they're in the buyer's own vocabulary. If your customers keep saying "they explained everything without making me feel dumb," the model can extract "good for first-timers" — an attribute you might never have thought to claim about yourself.
- **Your own site, especially the specific and concrete parts.** Your pages matter, but generic self-description barely moves the needle. "We provide quality service" gives the model nothing to extract. "We hand-set every stone and turn custom rings around in three weeks" gives it two clean, quotable attributes.
- **Comparisons and roundups.** When a third party puts you in a "best X for Y" list, they hand the model a ready-made attribute — the "for Y" part. Being in the right lists, framed the right way, is one of the most direct ways to shape what you're known for.

The through-line is that attributes are earned through corroboration and concreteness, not asserted through marketing copy. The model is pattern-matching on what the world says about you, weighted toward what many independent sources agree on.

## The most common attribute failures

When I look at a business that's getting cited but not converting through AI, I usually find one of three attribute problems.

**The generic blur.** The model can't find anything distinctive to say, so it defaults to bland, interchangeable descriptors — "a service provider in the area," "one option to consider." This is what happens when your entire footprint is safe, vague, feature-list marketing. You're not mischaracterized; you're just forgettable, and forgettable loses to specific every time.

**The stale attribute.** You've evolved — moved upmarket, added a specialty, changed your whole model — but the web still describes the old you. The model faithfully reports what it finds, which is your past. If you repositioned two years ago but your reviews, directory entries and old articles still describe the previous version, that's who the AI thinks you are.

**The wrong-magnet attribute.** You're characterized clearly, but as something that attracts the wrong customer. "The cheap option" when you're trying to sell premium. "Good for simple jobs" when you want the complex, high-margin work. This one's dangerous because it looks like success — you're getting named and described — while quietly routing your best-fit buyers elsewhere.

## Why the attribute layer decides the sale

It's worth sitting with why this matters more than raw mentions. In an AI-mediated purchase, the buyer often never visits your site before forming an opinion. They ask, the model characterizes a handful of options in a sentence each, and the buyer forms a shortlist from those sentences. Your carefully built website, your real strengths, your whole pitch — none of it gets a turn if the model's one-line summary of you didn't earn a click. The attribute layer is the entire first impression, and increasingly it's the only impression that happens before a shortlist forms.

That compresses a lot of pressure into very few words. It also means the attribute you're assigned tends to compound. If the model characterizes you as "best for beginners," you attract beginners, who then leave reviews describing a great beginner experience, which reinforces the attribute, which brings more beginners. The loop runs in whatever direction it started — toward the customer you want or the one you don't. That's exactly why the wrong-magnet attribute is so costly: it's not a one-time miss, it's a self-reinforcing groove that gets deeper every quarter you leave it alone.

The flip side is the opportunity. Because attributes compound, a deliberate push to establish the right one pays off past the initial effort — you're not just changing a description, you're changing which customers arrive and therefore which language accumulates about you next. Get the groove pointed the right way and it works for you on its own.

## How to shape what you're known for

You can't dictate your attributes, but you can absolutely influence the consensus the model reads. The work is deliberate and it's mostly about concreteness and repetition.

Start by deciding, honestly, the two or three things you actually want to be known for — and make sure they're true and defensible. This isn't a slogan exercise; the model will fact-check you against reality through everyone else's descriptions. Pick attributes you can back up.

Then make those attributes concrete and repeatable everywhere you control. Replace "quality service" with the specific, quotable proof of it. Say the same distinctive things consistently across your site, your profiles and your descriptions, in language a model can lift cleanly. Consistency across surfaces is itself a signal.

Next, get your customers and third parties saying it too, because that's the part the model weights most. Encourage reviews that speak to the attribute you want — not by scripting people, but by doing the thing so well and so visibly that they naturally describe it — and pursue placement in the roundups and comparisons where the "best for X" framing puts you in the right lane. When your own language and independent language converge on the same attributes, extraction gets easy and reliable.

Finally, audit it the same way you'd audit citations: ask the major AI engines to describe your business and to say who's best for specific use cases, and read the adjectives they hand back. That's your report card. If the words are generic, stale or pointed at the wrong customer, you've found the gap — and now you know exactly which consensus you need to go rebuild. My read is that in an AI-mediated market, being known for the right specific thing beats being known generally, and it's one of the few parts of AI visibility you can meaningfully steer.

### Key takeaways

- ♦ AI doesn't just decide whether to name you — it decides what to say about you, attaching attributes like 'affordable,' 'premium,' or 'best for beginners' that route buyers toward or away from you.
- ♦ Being named as the wrong thing can lose you the sale even when you appear in the same answer as a competitor — characterization often matters more than the citation itself.
- ♦ You don't declare your attributes; the model infers them from a consensus across your site, reviews, comparisons and forums — weighted heavily toward what many independent sources repeat.
- ♦ The three common failures are the generic blur (nothing distinctive to say), the stale attribute (the web describes an old version of you), and the wrong-magnet attribute (clear but attracting the wrong customer).
- ♦ Shape your attributes by making them concrete and repeatable everywhere you control, then getting customers and third parties to independently say the same true, specific things.
- ♦ Audit what you're known for by asking the major AI engines to describe you and to name who's best for specific use cases — the adjectives they return are your report card.

# Frequently asked questions

## Can I just tell AI what my business is known for?

Not directly — the model infers your attributes rather than accepting your claims. It reads a consensus of what the web says about you: your own pages, but also reviews, third-party descriptions, comparisons and forums, weighted toward what many independent sources agree on. You influence that consensus by making your distinctive attributes concrete and repeatable in the places you control, and by getting customers and credible third parties to independently describe you the same way. Marketing copy alone barely moves it; corroboration does.

## Why does AI describe my business with outdated information?

Because it faithfully reports the consensus it finds, and if you've repositioned, the web often still describes the previous version of you. Reviews, directory entries, old articles and comparison pages that predate your change keep feeding the model the old attributes. The fix is to refresh those surfaces — update your own descriptions to be specific and current, and generate fresh third-party signals (recent reviews, updated profiles, new roundup placements) that describe who you are now, so the recent consensus outweighs the stale one.

## How do I find out what AI thinks my business is known for?

Ask the major AI engines directly to describe your business and to say who's best for specific use cases in your category, then read the adjectives and framing they return. That's your report card. If the descriptions are generic ('one option to consider'), stale (describing an old version of you), or pointed at the wrong customer ('the cheap choice' when you sell premium), you've found the exact attribute gap to work on — and which consensus you need to rebuild through concrete, corroborated signals.

**About the author.** Scott Tischler is the Founder & Chairman of Alrecommend.ai and a practitioner-authority on AI search and Answer Engine Optimization. With 20+ years in marketing technology — including American Express, MetLife, and UBS — and executive and professional study at Wharton, Harvard, and Oxford, he helps businesses become the ones AI recommends.